

Temporal Measurement Instability in LLM Benchmarks: A Large-Scale Psychometric Audit of 5,227 Models

Anonymous ACL submission

Abstract

LLM benchmark scores are compared across model generations, yet whether individual items measure the same construct over time is rarely tested. Existing contamination detectors validate against web-overlap labels that conflate global exposure with differential functioning. We apply differential item functioning (DIF) analysis from educational measurement to six METABENCH benchmarks (5,227 models, 27,125 items), validated through a semi-synthetic injection framework with cross-subject matching. Between 2023 and 2024 model cohorts, 17.9–31.3% of items exhibit large temporal DIF; for MMLU, three convergent non-independence corrections yield approximately 27–29% (conservative bound: 13.5%; cluster bootstrap 95% CI [26.7%, 53.2%]; random-split baseline 0.18%). Web-overlap contamination labels are near-independent of DIF flags (Jaccard = 0.206), exposing a category error in prevailing validation methodology. This instability does not threaten aggregate rankings ($\rho \approx 0.997$) but reveals item-level temporal change, with 2024 models gaining on hard items while eroding on easy ones. LLM benchmarks need psychometric monitoring, not just score reporting.

1 Introduction

When we say “Model B outperforms Model A on MMLU,” we implicitly assume that the benchmark measures the same construct in the same way for both models. This assumption of temporal comparability underpins every cross-generation comparison published on leaderboards and in model reports, yet it is rarely tested. We use *measurement invariance* to denote the property that test items function identically across groups (Meredith, 1993), *differential item functioning* (DIF) as the item-level statistical test for violations, and *temporal measurement instability* as the empirical phenomenon when DIF is detected across time-defined

cohorts—a descriptive term denoting systematic changes in measurement properties, without implying degradation or attributing cause. Applying Mantel-Haenszel DIF analysis (Holland and Thayer, 1988) with ETS severity classification to the METABENCH response matrix (Kipnis et al., 2024),¹ we find that 31% of MMLU (Hendrycks et al., 2021) items receive ETS Category C classification ($|\Delta_{MH}| \geq 1.5$, $p < 0.05$) between 2023 and 2024 model cohorts—versus 0.18% under a random-split control. This instability does not threaten benchmark validity ($\rho \approx 0.997$ between full and stable-item rankings); it is diagnostic, revealing which capability dimensions shift across model generations.

Contamination detectors—from membership inference (Shi et al., 2024) to performance-based methods (Dekoninck et al., 2024; Liang, 2026; Tao et al., 2026)—validate primarily against web-overlap labels (Li et al., 2024b) that measure *global exposure*: whether an item exists in publicly accessible training corpora. Such labels do not capture *differential item functioning*: whether an item behaves differently across cohorts with varying exposure or training methodology (Figure 1). Direct comparison confirms near-independence (Jaccard = 0.206, $\phi = 0.012$), consistent across contamination subtypes.

We apply MH-DIF to the full METABENCH response matrix (5,227 models $\times$ 12,508 MMLU items), to our knowledge the largest-scale DIF analysis on LLM benchmarks, complementing concurrent work on anchor-item calibration (Habba et al., 2026) and addressing the fragility of existing detectors under post-training perturbation (Wang et al., 2026). Two structural controls, a sensitivity check, and replication across all six METABENCH benchmarks confirm that the temporal DIF rate (≈ 27 –

¹Analysis code and DIF audit scripts are available at <https://anonymous.4open.science/r/dif-audit>.




Figure 1: Schematic illustration of global vs. differential contamination via Item Characteristic Curves (ICCs). **Left:** global contamination produces no DIF; both cohorts encounter the item, and their ICCs overlap. **Right:** differential contamination produces DIF; the focal cohort’s ICC shifts because only that cohort was disproportionately exposed during training. The enrichment ratio of ≈ 1.0 between DIF flags and web-overlap labels is expected because most MMLU items fall into the left category.

29% after non-independence corrections) is not a methodological artifact (Sections 4.2, 4.4 and 4.7).

DIF should be understood as a diagnostic lens rather than a quality filter: DIF-flagged items concentrate discriminative signal ($z = -8$ to -49 vs. random removal; Section 4.6), and model families show distinctive DIF profiles.

Our contributions are methodological and empirical, not algorithmic:

1. **Ground truth meta-critique:** Web-overlap contamination labels (Li et al., 2024b) do not provide appropriate ground truth for item-level differential detectors. Direct empirical comparison on 9,229 items shows near-independence (Jaccard = 0.206, $\phi = 0.012$), exposing a validation methodology gap (Section 4.3).
2. **Semi-synthetic validation framework:** Cross-subject matching provides theoretical $\Delta\text{FPR} = 0$ guarantees for multi-domain benchmarks, with family-level clustering for non-independent models and per-item design-effect correction for correlated populations; replicated on GSM8K (Sections 4.5 and 4.7).
3. **Temporal measurement instability at scale:** Approximately 27–29% of MMLU items exhibit differential functioning across model generations after three convergent non-independence corrections; replication across all six METABENCH benchmarks yields $C\% = 17.9\text{--}31.3\%$, with knowledge/reasoning

benchmarks showing higher DIF rates than procedural ones (Sections 4.2, 4.6 and 4.7).

2 Related Work

Contamination Detection. A growing family of methods detects benchmark contamination through model behavior or training data analysis. Membership inference approaches estimate whether a model has seen specific items during training: Min-K% (Shi et al., 2024) and Min-K%++ (Zhang et al., 2024) use token-level log-likelihood statistics, while Time Travel (Golchin and Surdeanu, 2024) probes models with temporal cues and Oren et al. (2024) prove contamination through exchangeability tests. Performance-based detectors identify contamination through statistical anomalies (Dekoninck et al., 2024; Liang, 2026; Tao et al., 2026; Deng et al., 2024). Balloccu et al. (2024) provide a broader taxonomy of leakage detection strategies. These methods share a common validation practice: they verify detection accuracy against web-overlap or n-gram labels (Li et al., 2024b) that classify items based on verbatim matches in publicly accessible corpora. Such labels measure *global exposure* (whether an item exists in training data) rather than *differential item functioning* (whether an item behaves differently across model cohorts). This global-vs-differential distinction motivates our ground truth analysis (Section 4.3). Work on LLM memorization (Carlini et al., 2023; Tirumala et al., 2022) establishes that memorization depends on data frequency and model scale, reinforcing the distinction between global exposure and differential behavior. Wang et al. (2026) show that existing contamination detectors are fragile under reinforcement learning, further motivating behavioral approaches grounded in psychometric theory.

IRT in LLM Evaluation. IRT adoption in LLM evaluation has accelerated: Lalor et al. (2019) introduced IRT to NLP, Vania et al. (2021) assessed test set discriminability, and Rodriguez et al. (2021) demonstrated that not all evaluation examples are equally informative. Recent extensions include PSN-IRT (Zhou et al., 2026) for benchmark quality at scale, TinyBenchmarks (Maia Polo et al., 2024) for IRT-based item selection, scalability coefficients for detecting bad benchmark items (Hardy et al., 2026), IRT-based evaluation quality (Kraus and Arora, 2024), ATLAS (Li et al., 2025) for computerized adaptive testing of LLMs,

and competency-focused IRT evaluation in medical domains (Luo et al., 2025). The closest concurrent work is Growing Pains (Habba et al., 2026), which uses IRT calibration to establish anchor items for cross-temporal score equating, providing a prescriptive framework for maintaining comparability via fixed-parameter calibration (FPC; Kim and Kolen, 2019). Our work supplies the diagnostic foundation: identifying which items exhibit differential functioning and quantifying the magnitude of temporal instability. Their anchor-item approach can directly consume our stable-item set as pre-screened equating anchors (Kolen and Brennan, 2014). Our finding that 31% of items exhibit temporal DIF directly informs anchor selection by identifying which items should *not* serve as anchors. To our knowledge, no prior work applies DIF to benchmark temporal comparability—contamination detection and psychometric evaluation optimization have developed as separate threads.

DIF Methodology and Measurement Invariance.

The Mantel-Haenszel procedure for DIF detection was introduced by Holland and Thayer (1988) and remains the most widely used method in educational testing due to its simplicity and well-characterized statistical properties (Clauser and Mazor, 1998). Alternative approaches include logistic regression DIF (Swaminathan and Rogers, 1990), SIBTEST (Shealy and Stout, 1993), Lord’s χ^2 test (Lord, 1980), and the Raju area method (Raju, 1988); Millsap and Everson (1993), Zumbo (1999), and Penfield and Camilli (2007) provide comprehensive reviews. DIF analysis is a specific instance of measurement invariance testing (Meredith, 1993; Vandenberg and Lance, 2000; Millsap, 2011), which examines whether a measurement instrument functions equivalently across populations. MIMIC models offer a latent-variable approach to the same problem (Muthén, 1985; Woods, 2009), and Teresi and Fleishman (2007) reviews DIF methods across applied domains. Cross-subject matching—computing the stratification variable from domains other than the one being tested—has psychometric precedent in the DIF literature (Nandakumar, 1993; Donoghue and Allen, 1993); our adaptation addresses LLM-specific challenges of multi-domain local dependence and simultaneous contamination control at scale. We do not modify the MH procedure; our contribution is the application context and the empirical findings it reveals. The 5,227-examinee scale exceeds typical DIF

applications (Belzak, 2020). The impact-versus-DIF confound is addressed through matched-ability quintile analysis, IRT θ -conditioning (as sensitivity check), and a random-split control (Section 4.2).

Benchmark Quality and Design. Several concurrent approaches address benchmark reliability from complementary angles. The Benchmark Health Index (Zhu et al., 2026) and Benchmark² (Qian and Huang, 2026) perform meta-evaluation of static benchmark quality at a single snapshot; our temporal DIF analysis diagnoses *when* measurement properties shift. MMLU-CF (Zhao et al., 2025) reconstructs contamination-free items; DIF monitoring provides a lower-cost alternative when reconstruction is infeasible, identifying which existing items remain measurement-invariant rather than replacing them. Combined with the prescriptive equating framework of Growing Pains (Habba et al., 2026) (discussed above), these approaches span the full lifecycle from diagnosis (our work) to reconstruction (Zhao et al., 2025) to calibration (Habba et al., 2026). Maeda (2025) predict which item features associate with DIF in educational settings; we instead detect DIF empirically across LLM cohorts at scale. Dynamic benchmarks (Kiela et al., 2021; Li et al., 2024a; White, 2025) and self-evolving update frameworks (Ying et al., 2024; Xu et al., 2025) mitigate contamination by construction, yet none provides psychometric equating across versions. Broader surveys (Chang et al., 2024; Burnell et al., 2023) identify benchmark saturation and item-level opacity as systemic challenges that our DIF analysis directly addresses.

3 Method

We adopt Differential Item Functioning (DIF) analysis from educational measurement to assess whether individual benchmark items behave consistently across groups of LLMs. The procedure requires only a binary response matrix and group labels, with no access to model internals or training data. Below we define the data, the statistical test, the cohort designs, and a semi-synthetic validation framework.

3.1 Data

The primary dataset is the METABENCH response matrix (Kipnis et al., 2024), which records binary outcomes (1 = correct, 0 = incorrect) for 5,227 models evaluated on 12,508 MMLU items (Hendrycks

et al., 2021). Models span release dates from 2023 to 2024, drawn from the Open LLM Leaderboard. MMLU covers 57 subjects grouped into five broad domains (STEM, Humanities, Social Sciences, Other, and a mixed category). We supplement the response matrix with the web-overlap contamination labels of Li et al. (2024b), which classify 9,229 of 12,508 items (73.8% coverage) as “contaminated” or “clean” based on verbatim matches in publicly accessible web corpora. These labels serve as an external reference for evaluating the relationship between DIF flags and existing contamination indicators.

3.2 Mantel-Haenszel DIF and ETS Classification

In educational testing, Differential Item Functioning (DIF) occurs when an item’s statistical relationship to the latent trait θ differs between two groups after conditioning on ability level. An item with DIF is not merely harder for one group; it functions differently relative to what the group’s overall ability would predict. The standard terminology refers to test-takers as “examinees” and questions as “items.” In our setting, examinees are LLMs and items are benchmark questions.

The Mantel-Haenszel (MH) procedure (Holland and Thayer, 1988; Magis et al., 2010) detects DIF by stratifying examinees into K score bins based on total score (a proxy for θ) and computing a common odds ratio across strata. We use $K = 5$ equiprobable strata (quintile-based binning) following ETS recommended practice for large-scale assessments (Donoghue and Allen, 1993; Clauser and Mazor, 1998). Within each stratum k , the responses of the reference group and the focal group form a 2×2 contingency table:

$$\alpha_{\text{MH}} = \frac{\sum_{k=1}^K A_k D_k / T_k}{\sum_{k=1}^K B_k C_k / T_k}, \quad (1)$$

where A_k and B_k are the number of correct and incorrect responses in the reference group for stratum k , C_k and D_k are the corresponding counts for the focal group, and T_k is the total number of examinees in stratum k . Under the null hypothesis of no DIF, $\alpha_{\text{MH}} = 1$. The log odds ratio is converted to the ETS delta scale:

$$\Delta_{\text{MH}} = -2.35 \ln(\alpha_{\text{MH}}). \quad (2)$$

The Educational Testing Service (ETS) classifies DIF severity into three categories based on $|\Delta_{\text{MH}}|$. Category A (negligible) requires $|\Delta_{\text{MH}}| < 1.0$. Category C (large) requires both $|\Delta_{\text{MH}}| \geq 1.5$ and statistical significance at $p < 0.05$ from the MH chi-squared test. Category B (moderate) includes all remaining items: those with $1.0 \leq |\Delta_{\text{MH}}| < 1.5$, or $|\Delta_{\text{MH}}| \geq 1.5$ without statistical significance. The dual threshold in ETS C (effect size and significance) guards against flagging items that show large but unreliable effect sizes. Throughout this paper, we report the percentage of items classified as ETS C, denoted C%. At $N = 5,227$, significance ($p < 0.05$) is nearly automatically satisfied for any $|\Delta_{\text{MH}}| \geq 1.5$: only 3 of 3,918 C items lack significance, so the dual criterion effectively reduces to a single effect-size threshold; per-item DEFF correction (Section C) restores p -value selectivity by reducing effective N : under simple-mean DEFF, 269 items (6.9%) lose significance while retaining $|\Delta_{\text{MH}}| \geq 1.5$; 74% of these are near-miss ($p_{\text{adj}} \in [0.05, 0.10]$).

The choice of matching variable (the total score used for stratification) affects both validity and interpretation. We distinguish three matching strategies: (1) *standard matching* uses the within-benchmark total score on all items of the same benchmark; (2) *cross-subject matching*, adapting purified subtest matching (Swaminathan and Rogers, 1990) to the LLM multi-subject setting, uses the total score on all subjects *other than* the one containing the item under test, preventing the matching variable from absorbing injected contamination signals (Section 3.4); and (3) *external matching* uses the total score on a different benchmark entirely (e.g., MMLU total score as the matching variable when testing DIF on GSM8K items), providing an independent ability proxy uncontaminated by the target benchmark. Cross-subject and external matching both address matching variable contamination but at different granularities: cross-subject operates within a multi-domain benchmark, while external operates across benchmarks. Total score matching is sufficient: IRT θ correlates at $r = 0.9986$ with total scores (Section 4.4).

LLM-specific challenges. Applying MH-DIF to LLMs introduces assumptions absent in educational testing. First, overlapping training corpora violate local independence; permutation simulation

confirms LD-attributable FPR inflation < 0.001 pp under cross-subject matching (Section A). Second, fine-tuned variants and merges within model families are not independent examinees; graduated deduplication shows family structure does not inflate C% (Section 4.4). Third, bimodal ability distributions (base vs. instruction-tuned) motivate quintile-based rather than parametric stratification. Fourth, MMLU’s 57-subject structure raises a multidimensionality concern: if 2023 and 2024 cohorts differ on secondary ability dimensions, total-score matching may produce spurious DIF (Ackerman, 1992; Roussos and Stout, 1996). We address this through three lines of evidence: (i) domain-level ability profiles are nearly flat across cohorts (Cohen’s d : STEM = 0.26, Humanities = 0.16, Social Science = 0.15, Other = 0.15; range = 0.12), indicating uniform improvement that total-score matching adequately captures; (ii) cross-subject and within-subject matching show strong agreement ($\kappa = 0.80$, 5 subjects, 834 items; the 20% discordance is distributed across domains rather than concentrated in specific subjects), with cross-subject matching yielding 2.7 pp more C items—an upper bound on any multidimensional confound; (iii) cross-subject-unique C items do not concentrate in subjects with large cohort ability shifts, further dissociating multidimensional confound from the observed DIF signal (Sections B and 4.4). Unlike standard purified MH that conditions on total test score, we match on cross-subject total score to avoid criterion contamination when target-subject items themselves may be compromised—an adaptation specific to multi-domain LLM benchmarks where subjects may be differentially affected.

3.3 Cohort Designs

We construct five cohort comparisons: (1) *Temporal*: 2023 ($N = 1,080$) vs. 2024 ($N = 807$) models, the primary comparison; (2) *Ability-quintile*: Q1–Q2 vs. Q4–Q5 within the same temporal window; (3) *Within-family*: Llama-2 ($N = 161$) vs. Llama-3 ($N = 308$), isolating generational changes within a single family; (4) *Within-subject*: per-subject DIF to test domain concentration; and (5) *Random split*: two random halves of the 2023 cohort ($N = 540$ each) as a sanity check.

3.4 Semi-Synthetic Injection Framework

To evaluate whether MH-DIF can detect differential contamination when it is present, we construct a semi-synthetic framework with controlled ground

truth.

Procedure. From items classified as ETS A in the uninjected data (negligible DIF), we select 20% uniformly at random as the “contaminated” set. For each contaminated item, we flip incorrect responses to correct in the focal group at exploitation rate $p \in \{0.0, 0.3, 0.5, 0.7, 1.0\}$. We then run MH-DIF on the modified response matrix and measure detection performance. We report $\Delta\text{TPR} = \text{TPR}(\text{injected}) - \text{TPR}(\text{baseline})$ and $\Delta\text{FPR} = \text{FPR}(\text{injected}) - \text{FPR}(\text{baseline})$, where TPR is the true positive rate on injected items and FPR is the false positive rate on uninjected items. Each injection level is repeated across 5 random seeds for uncertainty estimation. Because binary response flipping represents a worst-case uniform injection pattern, reported ΔTPR values constitute upper-bound estimates; realistic contamination (partial memorization, format-specific advantages) would produce weaker signals.

Cross-subject matching. The choice of matching variable is critical for valid semi-synthetic evaluation. When contamination is injected into items within subject X , the total score on subject X absorbs part of the injection signal. Using this contaminated total score as the MH matching variable violates the conditional independence assumption and inflates FPR. Cross-subject matching resolves this by computing the matching variable from all subjects *other than* the one being tested.

Remark (Zero FPR under single-subject matching). Let $M_i^{(-s)} = \sum_{j \neq s} Y_{ij}$ be the matching variable for examinee i on focal subject s . Under single-subject injection (contamination confined to subject s), the tested item belongs to s and does not contribute to $M_i^{(-s)}$; injecting contamination into s leaves $M_i^{(-s)}$ unchanged. Therefore $\Delta\text{FPR} = 0$ by construction.

When contamination spans multiple subjects simultaneously, $M_i^{(-s)}$ may absorb signal from other contaminated subjects, potentially inflating FPR. Empirically, this guarantee extends to multi-domain contamination affecting $\leq 9\%$ of subjects ($k \leq 5$), with graceful degradation beyond this boundary (Section 4.5).

Empirically, $\Delta\text{FPR} = 0.000$ across all five MMLU domains (bootstrap 95% CI = $[0.000, 0.000]$). $\Delta\text{FPR} = 0$ is a constructive guarantee of the matching strategy, independent of injection realism, whereas ΔTPR constitutes

an upper-bound estimate under uniform injection. ΔFPR remains negligible for single-domain contamination ($k \leq 5$); when contamination spans many domains, within-subject matching should be preferred (Section 4.5). We adopt cross-subject matching as the standard protocol.

Following DIF convention (Wainer and Braun, 1993), we do not apply family-wise error correction: the dual ETS C threshold ($|\Delta_{\text{MH}}| \geq 1.5$ and $p < 0.05$) guards against false positives; BH-FDR at $q = 0.05$ reclassifies only 3 of 3,918 C items (Section A). Diagnostic procedures (Yen’s Q3, domain-specificity, DIF decomposition) and further details appear in Section A.

4 Experiments

We apply MH-DIF to the METABENCH response matrix (Kipnis et al., 2024) (5,227 models $\times$ 12,508 MMLU items (Hendrycks et al., 2021)), with temporal cohorts defined by release date (2023 vs. 2024). Cross-benchmark validation across all six METABENCH benchmarks appears in Section 4.7.

4.1 Validation: DIF Detection Power

As approximate parametric bounds, a Monte Carlo power analysis (1,000 replications per condition) using empirical 2PL parameters from PSN-IRT (Zhou et al., 2026) ($\bar{a} = 1.2$, $\bar{b} = -0.84$; 14,042 MMLU items) confirms that at $N \geq 80$ per cohort, ETS C achieves TPR of 0.68–0.80, FPR of 0.036–0.042, and AUC of 0.93–0.94 for items with $|\Delta_{\text{MH}}| \geq 1.4$ (details in Section A). Our temporal cohorts ($N_{\text{ref}} = 1,080$, $N_{\text{focal}} = 807$) far exceed this threshold. Power varies substantially with item discrimination: TPR ranges from 0.40 at $a = 0.5$ to 0.95 at $a = 2.0$; the gap between TPR at mean a (0.87) and the aggregate TPR (0.68) quantifies the power cost of parameter heterogeneity, consistent with Jensen’s inequality (Section A.2).

4.2 Temporal Comparability: 31% Measurement Instability

We compare 2023 ($N = 1,080$) against 2024 ($N = 807$) models; 31.3% of items receive ETS C classification ($|\Delta_{\text{MH}}| \geq 1.5$; 95% bootstrap CI: [29.4%, 38.1%]; family-level cluster bootstrap CI [26.7%, 53.2%]). Three independent non-independence corrections converge on approximately 27–29%: simple-mean DEFF correction (29.2%), graduated family de-duplication

at $K=20$ (29.1%; Table 7), and family-level cluster bootstrap lower bound (26.7%). The weighted-mean DEFF variant yields 13.5%, but its effective cluster size ($\bar{m}=254$) implies each family member is functionally identical—unrealistic for models with distinct training data; with only ~ 30 families, weighted-mean ICC estimates are unstable (Valliant et al., 2013). All reclassifications are significance-driven ($|\Delta_{\text{MH}}| \geq 1.5$ unchanged); the [13.5%, 29.2%] range reflects power uncertainty, not disagreement about DIF magnitude. We adopt ≈ 27 –29% as our primary corrected estimate, with 13.5% as a conservative lower bound (Sections C and D). Even the most conservative bound exceeds the random-split baseline by $75\times$. Leave-one-family-out analysis confirms no single family drives the signal ($|\Delta\text{C\%}| \leq 2.93$ pp). C% is stable across stratification choices ($K \geq 5$: 31.3%–35.9%; $K=5$ is most conservative); the non-monotonic $K=10$ peak (35.9%) reflects finer strata capturing ability–DIF interactions, $K=20$ instability from small strata (~ 94 models), and $K=3$ inflates to 39.4% from under-stratification.

Three controls confirm genuine temporal change (Table 1): random splits yield C% of 0.18% (2023) and 1.3% (2024), both $\geq 23\times$ below temporal C%; ability-quintile comparison yields 2.9%; and 1:1 nearest-neighbor matching on total score ($N=638/\text{group}$, Cohen’s $d = -0.17$) yields C% = $28.9\% \pm 0.1\%$, confirming at most ~ 2.4 pp is attributable to ability mismatch ($p < 0.01$; Section A).

Across all 57 subjects, C% ranges from 20.7% to 48.4% (mean $\approx 30\%$; Table 1). DIF rates are domain-invariant (Kruskal-Wallis $p = 0.78$), but DIF *direction* exhibits domain dependence: STEM items favor 2024 models (55.4%), Humanities favor 2023 (64.4%). Domain-level accuracy gains explain $< 0.1\%$ of item-level DIF variance (pseudo $R^2 = 0.0007$); a Simpson’s paradox confirms DIF is associated with item properties rather than aggregate gains (OR reverses from 1.55 to 0.73 after controlling for difficulty; Section B). A difficulty–direction interaction emerges ($\rho = -0.597$): 2024 models gain on hard items while eroding on easy ones. The conclusion is robust across thresholds (18.3% at $|\Delta_{\text{MH}}| \geq 2.0$, 10.2% at ≥ 2.5).

4.3 Ground Truth Mismatch

Direct comparison of DIF flags against the web-overlap labels of Li et al. (2024b) on 9,229 annotated items confirms that the two methods cap-

Table 1: DIF landscape across cohort definitions. ETS C denotes the percentage of items with $|\Delta_{MH}| \geq 1.5$. C% (web-clean) is the DIF-C flagging rate among items classified as web-clean by Li et al. (2024b). Enrichment is the ratio of observed to expected overlap between DIF flags and web-overlap labels. Dashes (–) indicate comparisons where Li et al. labels are not applicable.

Cohort Definition	N_{ref} / N_{focal}	C%	C% (web-clean)	Enrich.
Temporal (2023 vs. 2024)	1,080 / 807	31.3	0.318	0.964
Ability quintile (Q1–2 vs. Q4–5)	~2,600 / ~2,600	2.9	0.027	1.19
Within-family (Llama-2 vs. 3)	161 / 308	61.3	–	–
Random split 2023 (sanity)	540 / 540	0.18	–	–
Random split 2024 (sanity)	~400 / ~400	1.3	–	–
Within-subject (avg. 5 subj.)	varies	38.1	–	–

ture orthogonal signals. The enrichment ratio is 0.964 (bootstrap 95% CI: [0.94, 1.02]; Table 1): DIF-C rates are 30.7% among web-contaminated items and 31.8% among clean items. Association tests confirm near-independence ($\phi = 0.012$, OR = 1.05 [0.96, 1.15], Jaccard = 0.206, within permutation null [0.203, 0.222]; $p = 0.88$). This near-independence holds across contamination subtypes. The null result does not indicate detection failure: semi-synthetic injection produces a clear dose-response ($\Delta TPR = +0.508$ at $p = 0.5$; Figure 2). The disconnect arises because web-overlap labels measure *global exposure* (whether an item exists in training corpora), while DIF measures *differential functioning* (whether an item’s relationship to ability changes across cohorts). The Li et al. labels are the wrong ground truth for a differential method.

4.4 Structural Analysis

Two structural controls and a sensitivity check fail to reduce temporal C% below 31% (Table 2): matched-ability quintiles, within-subject DIF (C% = 38.1%; cross-subject C% is more conservative due to lower local dependence), and IRT θ -conditioning ($r(\theta, \text{total score}) = 0.9986$). Logistic regression on nine subjects (1,787 items; 9/57—full-benchmark rates require further analysis) detects an additional 17.7% with non-uniform DIF ($|\beta_{\text{interaction}}| \geq 0.5$; effect-size criterion; LR p -values not adjusted for clustering); the union reaches 47% (Jaccard = 0.28), confirming that MH-based ≈ 27 –29% captures only the uniform component. Ability-matched analysis yields C% = 28.9% (temporal) and 28.3% (Llama-2 vs. 3)—consistent matched values (~ 28 –29%) despite vastly different unmatched rates (Section B).




Figure 2: Semi-synthetic dose-response for DIF detection on MMLU. ΔTPR increases monotonically with exploitation rate p , confirming that MH-DIF is sensitive to differential item behavior when present. Baseline (no injection) produces $\Delta TPR \approx 0$.

Table 2: Structural controls and sensitivity checks for temporal DIF. None reduces C% below 31%; θ -conditioning serves as a sensitivity check confirming total-score adequacy.

Approach	Mechanism	Temp. C%	TPR ($p=0.5$)	Outcome
Baseline (ext.)	Cross-bench, total score	0.314 $\pm$.009	0.834 $\pm$.027	Reference
Matched-ability	Within-quintile temp. DIF	~ 0.31	–	Enrich. 0.94–1.04
θ -cond. (sens.)	θ replaces total score	0.312	0.876	$r(\theta, \text{total})=0.9986$
Within-subject	Per-subject DIF	–	–	C% = 38.1% (>25%)

4.5 Semi-Synthetic Framework Validation

Cross-subject matching achieves $\Delta FPR = 0.000$ across all five MMLU subjects (4/5 with $\Delta TPR > 0.3$ at $p = 0.5$, upper-bound estimate under uniform binary flipping; Section 3.4), while within-subject matching inflates FPR by +0.186 on average. The $\Delta FPR = 0$ guarantee holds for single-subject contamination ($k \leq 5$, $\Delta FPR < 0.007$). Under broader contamination spanning $k = 10$ subjects (18% of domains), ΔFPR rises to 0.047 ± 0.022 ; at $k = 20$, to 0.069 ± 0.018 (Table 6 and fig. 3). Even worst-case multi-subject inflation (~ 7 pp) cannot account for the observed 31% C% rate, which exceeds the 0.18% random baseline by $> 170\times$. Under DEFF-adjusted effective sample sizes, worst-case ΔFPR at $k=20$ (0.069) represents up to 50% of the conservative 13.5% lower bound, underscoring that cross-subject matching should be used with $k \leq 5$ in practice.

Table 3: Temporal DIF across six METABENCH benchmarks. C% = ETS C items ($|\Delta_{MH}| \geq 1.5$) between 2023 and 2024 cohorts; ρ/τ = Spearman/Kendall correlation between full and stable-item rankings; Flip = pairwise rank-flip rate; Ratio = C%/Rand. C%, with baseline varying by benchmark. All six show substantial DIF (45–174 $\times$ above random baseline), with knowledge/reasoning benchmarks exhibiting higher C% than procedural ones.

Benchmark	Items	C%	ρ	τ	Flip%	Rand. C%	Ratio
MMLU	12,508	31.3	0.9965	0.9600	2.00	0.18	174 $\times$
WinoGrande	1,267	30.3	0.9943	0.9432	2.84	0.32	95 $\times$
TruthfulQA	817	29.7	0.9971	0.9586	2.07	0.37	80 $\times$
ARC	1,172	26.9	0.9962	0.9538	2.31	0.60	45 $\times$
GSM8K	1,319	18.9	0.9993	0.9843	0.78	0.30	63 $\times$
HellaSwag	10,042	17.9	0.9986	0.9722	1.39	0.28	64 $\times$

4.6 Downstream Impact

Stable-item ranks are preserved across all six benchmarks ($\rho = 0.997$, $\tau = 0.958$; flip rate $\leq 2.8\%$; Table 3). DIF-flagged items concentrate discriminative signal ($z = -8$ to -49 vs. random removal). Model families show distinctive profiles: Llama-3 derivatives gain +53 median rank positions, Mistral derivatives lose 35 (Figure 4). Among the 200 most displaced models (top 3.8%; full-sample flip rate = 2.0%), 33.6% of pairwise rankings reverse. Removing C-flagged items disproportionately disadvantages 2024 models (cohort gap = 96.2 positions, 40 $\times$ random; $p < 0.001$).

4.7 Cross-Benchmark Validation

Temporal DIF extends across all six METABENCH benchmarks (Table 3): C% ranges from 17.9% (HellaSwag (Zellers et al., 2019)) to 31.3% (MMLU), with knowledge/reasoning benchmarks consistently higher than procedural ones (C% = 26.9–31.3% vs. 17.9–18.9%). These differences may partly reflect item difficulty distributions; we treat them as descriptive.

Semi-synthetic injection on GSM8K confirms detection power ($\Delta\text{TPR} = +0.514$; Section E); exploratory GSM1K analysis ($r = -0.478$, $p = 0.005$; Figure 5) suggests DIF captures beyond memorization. DIF remains substantial within 6-month windows (C% = 34.35%).

5 Discussion

5.1 What Drives Temporal DIF?

Temporal DIF detects *that* items function differently across cohorts, not *why*. Three non-exclusive mechanisms may drive the observed ≈ 27 –29% rate: (1) differential contamination (asymmetric training data exposure); (2) capability evolution (genuine skill acquisition); (3) training methodology shifts (RLHF, post-training). The ETS baseline of 5–15% C% (Clauser and Mazor, 1998) is for concurrent demographic groups sharing the same test administration; the comparison to temporal

item-specific. *Item quality defects* are ruled out: expert-annotated errors (Gema et al., 2024) show no enrichment among DIF-C items (enrichment = 1.10, 95% CI [0.92, 1.27], $p = 0.23$). *Web-overlap labels* (Li et al., 2024b) are near-independent of DIF flags (Jaccard = 0.206, indistinguishable from permutation null; Section 4.3), a category error in ground truth; MMLU-CF (Zhao et al., 2025) complements our approach via item reformulation.

The difficulty–direction interaction ($\rho = -0.597$; Section 4.2) provides suggestive evidence for capability evolution: hardest-decile DIF-C items are 88.7% focal-favoring (2024 better), easiest-decile 96.2% reference-favoring. This pattern is more consistent with skill acquisition than contamination, which would target items regardless of difficulty; however, temporal DIF is necessary but not sufficient for construct-irrelevant change, and observational DIF alone cannot separate contamination from capability evolution (Suk and Lyu, 2026). The high aggregate $\rho \approx 0.997$ does not undermine DIF’s utility: aggregate stability masks subject-level instability (mean per-subject $\rho = 0.950$; 22/57 below 0.95; five below 0.90). DIF monitoring enables targeted replacement (Zhao et al., 2025; Habba et al., 2026).

5.2 Recommended DIF Audit Protocol

Item-level properties do not predict DIF susceptibility (AUROC = 0.60; Section G), so monitoring cannot be replaced by pre-screening. Circularity analysis confirms DIF removal is the *least* disruptive non-random strategy ($z = -9.5$ vs. accuracy-gap $z = -157$; Figure 4), motivating equating (Kolen and Brennan, 2014; Habba et al., 2026). We recommend MH-DIF audits whenever ≥ 500 models accumulate, with three action levels (Figure 6): *Monitor* (C% < 15%), *Flag* (15–25%: dual rankings), *Act* ($\geq 25\%$: stable-item rankings as primary). Subject-level monitoring is essential (mean $\rho = 0.950$; 22/57 subjects below 0.95).

6 Conclusion

Across six METABENCH benchmarks, 17.9–31.3% of items exhibit temporal measurement instability between 2023 and 2024 model cohorts (≈ 27 –29% for MMLU after non-independence corrections; conservative bound: 13.5%). Web-overlap labels are near-independent of DIF flags (Jaccard = 0.206); rankings are preserved ($\rho \approx 0.997$), motivating routine temporal DIF audits.

Limitations

Methodological scope. MH-DIF detects only uniform DIF; the 17.7% of items showing non-uniform patterns (ability-dependent DIF) may be underestimated by the Mantel-Haenszel statistic, and future work could employ logistic regression DIF or IRT-based methods for fuller characterization. Logistic regression reveals this additional non-uniform DIF (effect-size criterion; p -values not adjusted for within-family clustering); the union reaches 47% on 9/57 subjects, making the corrected ≈ 27 –29% conservative but subject to further validation at full-benchmark scale. More fundamentally, DIF detects *that* items function differently but cannot determine *why*—the signal may reflect contamination, capability change, or training methodology shifts. A definitive causal decomposition remains out of reach without access to model training data. The semi-synthetic framework uses binary response flipping (incorrect $\rightarrow$ correct) to simulate differential contamination, representing a worst-case uniform exploitation pattern. Real-world contamination involves partial memorization, reasoning shortcuts, and format-specific advantages; detection power under such realistic patterns is untested. The moderate $\Delta\text{TPR} = +0.508$ at $p = 0.5$ reflects sensitivity to severe differential behavior; the framework’s primary methodological contribution—cross-subject matching achieving $\Delta\text{FPR} = 0.000$ —is a matching strategy property independent of injection realism. Our uniform injection model is difficulty-agnostic, whereas real contamination likely correlates with item difficulty. Difficulty-dependent injection would concentrate DIF on easy items, potentially increasing detection power for that subset while reducing sensitivity to hard-item DIF; the observed difficulty–direction interaction ($\rho = -0.597$) suggests that actual patterns are more complex than uniform models capture.

Examinee non-independence. The MH χ^2 test assumes independent examinees, yet LLMs sharing base models, training data, or architectures violate this assumption. Empirical mitigations partially address this: permutation simulation bounds LD-attributable FPR inflation to <0.001 pp, and graduated de-duplication ($K=20$: $C\% = 29.1\%$) shows family structure does not inflate results. French and Finch (2010) show that MH maintains nominal size under moderate clustering in educational settings, consistent with our empirical findings; however, the

extreme relatedness of some LLM families (e.g., hundreds of Llama-2 fine-tunes) exceeds the clustering levels studied in that work. Cluster bootstrap resampling at the family level (29 families with ≥ 3 members, 1,882 models, $B = 1,000$) yields a 95% CI of [26.7%, 53.2%], $3.1\times$ wider than the model-level CI [29.4%, 38.1%], partly because LLM families are cohort-specific (no family spans both 2023 and 2024), so resampling families simultaneously varies the reference/focal composition. Despite this conservative bound, the lower bound far exceeds the random-split baseline (0.18%), confirming robustness to within-family dependence.

Matching boundary and sample size. Cross-subject matching assumes contamination affects a minority of domains; when contamination spans >10 subjects ($>18\%$ of MMLU), ΔFPR rises to 0.026–0.069. Model-level bootstrap may underestimate CI width due to within-family correlation; family-level cluster bootstrap (preceding paragraph) confirms robustness despite a $3.1\times$ wider CI. The GSM1K correlation analysis ($N = 33$) has limited statistical power and is exploratory.

Temporal granularity. Our temporal cohorts are defined by model release date, which may lag training data cutoff by months. This coarse granularity, combined with within-cohort heterogeneity (e.g., 2023-Q1 vs. 2024-Q4 spans a larger gap than 2023-Q4 vs. 2024-Q1), likely attenuates DIF estimates, making ≈ 27 –29% a possible underestimate of true temporal instability.

Generalizability. All six benchmarks use multiple-choice or binary scoring; generalizability to open-ended generation (e.g., AlpacaEval), coding (e.g., HumanEval), or partial-credit benchmarks is untested and would require polytomous DIF methods. Temporal cohorts are defined by calendar year, a coarse grouping; finer-grained definitions (e.g., by training data cutoff or alignment method) might reveal more targeted patterns. All models are drawn from the Open LLM Leaderboard; generalizability to closed-source models cannot be assessed.

References

Terry A Ackerman. 1992. A didactic explanation of item bias, item impact, and item validity from a multidimensional perspective. *Journal of Educational Measurement*, 29(1):67–91.

805	Simone Balloccu, Patrícia Schmidtova, Mateusz Lango,	Du, Mohammad Reza Ghasemi Madani, Claire Bar-	860
806	and Ondrej Dusek. 2024. Leak, cheat, repeat: Data	ale, Robert McHardy, Joshua Harris, Jean Kad-	861
807	contamination and evaluation malpractices in closed-	dour, Emile van Krieken, and Pasquale Minervini.	862
808	source LLMs . In <i>Proceedings of the 18th Conference</i>	2024. Are we done with MMLU? <i>Preprint</i> ,	863
809	<i>of the European Chapter of the Association for Com-</i>	arXiv:2406.04127.	864
810	<i>putational Linguistics (EACL)</i> .		
811	William Belzak. 2020. Testing differential item func-	Shahriar Golchin and Mihai Surdeanu. 2024. Time	865
812	tioning in small samples . <i>Multivariate Behavioral</i>	travel in LLMs: Tracing data contamination in large	866
813	<i>Research</i> .	language models . In <i>Proceedings of the International</i>	867
		<i>Conference on Learning Representations (ICLR)</i> .	868
814	Ryan Burnell, Wout Schellaert, John Burden, Tomer D.	Eliya Habba, Itay Itzhak, Asaf Yehudai, Yotam Per-	869
815	Ullman, Fernando Martinez-Plumed, Joshua B.	litz, Elron Bandel, Michal Shmueli-Scheuer, Leshem	870
816	Tenenbaum, Danaja Rutar, Lucy G. Cheke, Jascha	Choshen, and Gabriel Stanovsky. 2026. Grow-	871
817	Sohl-Dickstein, Melanie Mitchell, Douwe Kiela,	ing pains: Extensible and efficient LLM bench-	872
818	Murray Shanahan, Ellen M. Voorhees, Anthony G.	marking via fixed parameter calibration . <i>Preprint</i> ,	873
819	Cohn, Joel Z. Leibo, and Jose Hernandez-Orallo.	arXiv:2604.12843.	874
820	2023. Rethink reporting of evaluation results in AI		
821	Science .	Matthew Hardy, Juan Gilbert, and Benjamin Domingue.	875
		2026. Efficient detection of bad benchmark	876
822	Nicholas Carlini, Daphne Ippolito, Matthew Jagielski,	items with novel scalability coefficients . <i>Preprint</i> ,	877
823	Katherine Lee, Florian Tramèr, and Chiyuan Zhang.	arXiv:2603.24999.	878
824	2023. Quantifying memorization across neural lan-		
825	guage models . In <i>International Conference on Learn-</i>	Dan Hendrycks, Collin Burns, Steven Basart, Andy	879
826	<i>ing Representations (ICLR)</i> .	Zou, Mantas Mazeika, Dawn Song, and Jacob Stein-	880
		hardt. 2021. Measuring massive multitask language	881
827	Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu,	understanding . In <i>Proceedings of the International</i>	882
828	Kaijie Zhu, Hao Chen, Linyi Yang, Xiaoyuan Yi,	<i>Conference on Learning Representations (ICLR)</i> .	883
829	Cunxiang Wang, Yidong Wang, Wei Ye, Yue Zhang,		
830	Yi Chang, Philip S. Yu, Qiang Yang, and Xing Xie.	Paul W. Holland and Dorothy T. Thayer. 1988. Dif-	884
831	2024. A survey on evaluation of large language mod-	ferential item functioning and the Mantel-Haenszel	885
832	els . <i>ACM Transactions on Intelligent Systems and</i>	procedure. In Howard Wainer and Henry I. Braun,	886
833	<i>Technology</i> .	editors, <i>Test Validity</i> . Lawrence Erlbaum Associates.	887
834	Brian E. Clauser and Kathleen M. Mazor. 1998. Using	Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh	888
835	statistical procedures to identify differentially func-	Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vid-	889
836	tioning test items . <i>Educational Measurement: Issues</i>	gen, Grusha Prasad, Amanpreet Singh, Pratik Ring-	890
837	<i>and Practice</i> .	shia, Zhiyi Ma, Tristan Thrush, Sebastian Riedel,	891
		Zeeraq Talat, Pontus Stenetorp, and Robin Jia. 2021.	892
838	Jasper Dekoninck, Mark Niklas Mueller, and Martin	Dynabench: Rethinking benchmarking in NLP . In	893
839	Vechev. 2024. ConStat: Performance-based contam-	<i>Proceedings of the 2021 Conference of the North</i>	894
840	ination detection in large language models . In <i>Ad-</i>	<i>American Chapter of the Association for Computa-</i>	895
841	<i>vances in Neural Information Processing Systems</i>	<i>tional Linguistics (NAACL)</i> .	896
842	<i>(NeurIPS)</i> .		
843	Chunyan Deng, Yilun Zhao, Xiangru Tang, Mark Ger-	Seonghoon Kim and Michael Kolen. 2019. Application	897
844	stein, and Arman Cohan. 2024. Investigating data	of IRT fixed parameter calibration to multiple-group	898
845	contamination in modern benchmarks for large lan-	test data . <i>Applied Measurement in Education</i> .	899
846	guage models . In <i>Proceedings of the 2024 Confer-</i>		
847	<i>ence of the North American Chapter of the Associa-</i>	Alon Kipnis, Konstantinos Voudouris, Luca Maria	900
848	<i>tion for Computational Linguistics (NAACL)</i> .	Schulze Buschoff, and Eric Schulz. 2024. metabench	901
		– a sparse benchmark of reasoning and knowledge in	902
849	John R. Donoghue and Nancy L. Allen. 1993. An em-	large language models . <i>Preprint</i> , arXiv:2407.12844.	903
850	pirical examination of the Mantel-Haenszel statistic		
851	for the study of differential item functioning. <i>Applied</i>	Leslie Kish. 1965. <i>Survey Sampling</i> . John Wiley &	904
852	<i>Measurement in Education</i> , 6(2):163–179.	Sons, New York.	905
853	Brian F. French and W. Holmes Finch. 2010. Hierar-	Michael J. Kolen and Robert L. Brennan. 2014. <i>Test</i>	906
854	chical logistic regression: Accounting for multilevel	<i>Equating, Scaling, and Linking: Methods and Prac-</i>	907
855	data in DIF detection. <i>Journal of Educational Mea-</i>	<i>tices</i> , 3rd edition. Springer.	908
856	<i>surement</i> , 47(3):299–317.		
857	Aryo Pradipta Gema, Joshua Ong Jun Leang, Giwon	Batu El-Houl Kraus and Shashank Arora. 2024. Using	909
858	Hong, Alessio Devoto, Alberto Carlo Maria Man-	IRT to measure the quality of LLM evaluations. In	910
859	cino, Rohit Saxena, Xuanli He, Yu Zhao, Xiaotang	<i>Proceedings of the 2024 Conference on Empirical</i>	911
		<i>Methods in Natural Language Processing (EMNLP)</i> .	912

913	John Lalor, Hao Wu, and Hong Yu. 2019. Learning latent parameters without human response patterns: Item response theory with artificial crowds . In <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)</i> .	966
914		967
915		968
916		
917		969
918		970
919		971
		972
920	Peiyang Li, Xiaobo Tang, Sai Chen, Ying Cheng, Ronald Metoyer, Tianyi Hua, and Nitesh V. Chawla. 2025. ATLAS: Adaptive testing for LLM evaluation: A psychometric alternative to static benchmarks . Preprint, arXiv:2511.04689.	
921		
922		973
923		974
924		975
		976
		977
925	Yucheng Li, Frank Guerin, and Chenghua Lin. 2024a. LatestEval: Addressing data contamination in language model evaluation through dynamic and time-sensitive test construction . In <i>Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)</i> .	
926		978
927		979
928		980
929		
930	Yucheng Li, Yunhao Guo, Frank Guerin, and Chenghua Lin. 2024b. An open-source data contamination report for large language models . In <i>Findings of the Association for Computational Linguistics: EMNLP 2024</i> .	
931		
932		
933		
934		
935	Rongzhen Liang. 2026. DVD: A robust method for detecting variant contamination in large language model evaluation . Preprint, arXiv:2601.04895.	
936		
937		
938	Frederic M. Lord. 1980. <i>Applications of Item Response Theory to Practical Testing Problems</i> . Lawrence Erlbaum Associates, Hillsdale, NJ.	
939		
940		
941	Zhimeng Luo, Lixin Wu, Adam Frisch, and Daqing He. 2025. Measuring competency, not performance: Item-aware evaluation across medical benchmarks . Preprint, arXiv:2509.24186.	
942		
943		
944		
945	Hotaka Maeda. 2025. Finding words associated with DIF: Predicting DIF using LLMs and explainable AI . Preprint, arXiv:2502.07017.	
946		
947		
948	David Magis, Sébastien Béland, Francis Tuerlinckx, and Paul De Boeck. 2010. A general framework and an R package for the detection of dichotomous differential item functioning . <i>Behavior Research Methods</i> , 42(3):847–862.	
949		
950		
951		
952		
953	Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. tinyBenchmarks: evaluating LLMs with fewer examples . In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> .	
954		
955		
956		
957		
958	William Meredith. 1993. Measurement invariance, factor analysis and factorial invariance . <i>Psychometrika</i> .	
959		
960	Roger E. Millsap. 2011. Statistical Approaches to Measurement Invariance . Routledge, New York.	
961		
962	Roger E. Millsap and Howard T. Everson. 1993. Methodology review: Statistical approaches for assessing measurement bias . <i>Applied Psychological Measurement</i> , 17(4):297–334.	
963		
964		
965		
	Bengt Muthén. 1985. A method for studying the homogeneity of test items with respect to other relevant variables . <i>Journal of Educational Statistics</i> .	966
		967
		968
	Ratna Nandakumar. 1993. Simultaneous DIF amplification and cancellation: Shealy-Stout’s test for DIF. <i>Journal of Educational Measurement</i> , 30(4):293–311.	969
		970
		971
		972
	Yonatan Oren, Nicole Meister, Niladri S. Chatterji, Faisal Ladhak, and Tatsunori B. Hashimoto. 2024. Proving test set contamination in black-box language models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	973
		974
		975
		976
		977
	Randall D. Penfield and Gregory Camilli. 2007. Differential item functioning and item bias . In <i>Handbook of Statistics</i> , volume 26, pages 125–167. Elsevier.	978
		979
		980
	Qi Qian and Chengsong Huang. 2026. Benchmark²: Systematic evaluation of LLM benchmarks . Preprint, arXiv:2601.03986.	981
		982
		983
	Nambury S. Raju. 1988. The area between two item characteristic curves . <i>Psychometrika</i> .	984
		985
	J. N. K. Rao and A. J. Scott. 1992. A simple method for the analysis of clustered binary data. <i>Biometrics</i> , 48(2):577–585.	986
		987
		988
	Pedro Rodriguez, Joe Barrow, Alexander Miserlis Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. 2021. Evaluation examples are not equally informative: How should that change NLP leaderboards? In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)</i> , pages 4486–4503.	989
		990
		991
		992
		993
		994
		995
	Louis A Roussos and William F Stout. 1996. A new view of the effect of multidimensionality on the Mantel-Haenszel and SIBTEST procedures. <i>Journal of Educational Measurement</i> , 33(3):278–295.	996
		997
		998
		999
	Robin Shealy and William Stout. 1993. A model-based standardization approach that separates true bias/DIF from group ability differences and detects test bias/DTF as well as item bias/DIF . <i>Psychometrika</i> .	1000
		1001
		1002
		1003
		1004
	Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. 2024. Detecting pretraining data from large language models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	1005
		1006
		1007
		1008
		1009
		1010
	Youmi Suk and Weijun Lyu. 2026. Rethinking differential item functioning from a causal inference perspective. <i>Psychometrika</i> . Forthcoming.	1011
		1012
		1013
	Hariharan Swaminathan and H. Jane Rogers. 1990. Detecting differential item functioning using logistic regression procedures . <i>Journal of Educational Measurement</i> .	1014
		1015
		1016
		1017

1018	Yongding Tao, Tian Wang, Yihong Dong, Huanyu Liu,	Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali	1072
1019	Kechi Zhang, Xiaolong Hu, and Ge Li. 2026. Detect-	Farzaneh, and Yejin Choi. 2019. HellaSwag: Can a	1073
1020	ing data contamination from reinforcement learning	machine really finish your sentence? In <i>Proceedings</i>	1074
1021	post-training for large language models . In <i>Proceed-</i>	of the 57th Annual Meeting of the Association for	1075
1022	<i>ings of the International Conference on Learning</i>	<i>Computational Linguistics (ACL)</i> , pages 4791–4800.	1076
1023	<i>Representations (ICLR)</i> .		
1024	Jeanne A. Teresi and John A. Fleishman. 2007. Differen-	Jingyang Zhang, Jingwei Sun, Eric Yeats, Yang Ouyang,	1077
1025	tial item functioning and health assessment . <i>Quality</i>	Martin Kuo, Jianyi Zhang, Hao Yang, and Hai Li.	1078
1026	<i>of Life Research</i> .	2024. Min-K%++: Improved baseline for detect-	1079
1027	Kushal Tirumala, Aram Markosyan, Luke Zettlemoyer,	ing pre-training data from large language models .	1080
1028	and Armen Aghajanyan. 2022. Memorization with-	<i>Preprint</i> , arXiv:2404.02936.	1081
1029	out overfitting: Analyzing the training dynamics of	Qihao Zhao, Yangyu Huang, Tengchao Jing, Xiaom-	1082
1030	large language models . In <i>Advances in Neural Infor-</i>	ing Shi, Yanan Liang, Jiaqing Du, Wenbo He, and	1083
1031	<i>mation Processing Systems (NeurIPS)</i> .	Yankai Lin. 2025. MMLU-CF: A contamination-	1084
1032	Richard Valliant, Jill A. Dever, and Frauke Kreuter.	free multi-task language understanding benchmark .	1085
1033	2013. Practical Tools for Designing and Weighting	In <i>Proceedings of the 63rd Annual Meeting of the</i>	1086
1034	Survey Samples . Springer, New York.	<i>Association for Computational Linguistics (ACL)</i> .	1087
1035	Robert J. Vandenberg and Charles E. Lance. 2000. A	Hongli Zhou, Hui Huang, Ziqing Zhao, Lvyuan Han,	1088
1036	review and synthesis of the measurement invariance	Huicheng Wang, Kehai Chen, Muyun Yang, Wei	1089
1037	literature: Suggestions, practices, and recommen-	Bao, Jian Dong, Bing Xu, Conghui Zhu, Hailong	1090
1038	dations for organizational research . <i>Organizational</i>	Cao, and Tiejun Zhao. 2026. Lost in benchmarks?	1091
1039	<i>Research Methods</i> .	rethinking large language model benchmarking with	1092
1040	Clara Vania, Phu Mon Htut, William Huang, Dhara	item response theory . In <i>Proceedings of the AAAI</i>	1093
1041	Mungra, Richard Yuanzhe Pang, Jason Phang,	<i>Conference on Artificial Intelligence (AAAI)</i> .	1094
1042	Haokun Liu, Kyunghyun Cho, and Samuel R. Bow-	Longyuan Zhu, Hairan Hua, Linlin Miao, and Bing	1095
1043	man. 2021. Comparing test sets with item response	Zhao. 2026. Benchmark health index: A system-	1096
1044	theory . In <i>Proceedings of the 59th Annual Meeting of</i>	atic framework for benchmarking the benchmarks of	1097
1045	<i>the Association for Computational Linguistics (ACL)</i> .	LLMs . <i>Preprint</i> , arXiv:2602.11674.	1098
1046	Howard Wainer and Henry I. Braun. 1993. <i>Differential</i>	Bruno D. Zumbo. 1999. <i>A Handbook on the Theory and</i>	1099
1047	<i>Item Functioning</i> . Lawrence Erlbaum Associates,	<i>Methods of Differential Item Functioning (DIF): Lo-</i>	1100
1048	Hillsdale, NJ.	<i>gistic Regression Modeling as a Unitary Framework</i>	1101
1049	Haonan Wang, Haoyang Li, Boeun Ko, and Hongyang	<i>for Binary and Likert-Type (Ordinal) Item Scores</i> .	1102
1050	Zhang. 2026. On the fragility of benchmark contami-	Directorate of Human Resources Research and Eval-	1103
1051	nation detection in reasoning models . In <i>Proceedings</i>	uation, Department of National Defense, Ottawa.	1104
1052	<i>of the International Conference on Learning Repre-</i>		
1053	<i>sentations (ICLR)</i> .	A Method Details	1105
1054	Colin White. 2025. LiveBench: A challenging,	A.1 Diagnostics	1106
1055	contamination-free LLM benchmark . In <i>Proceed-</i>	Three diagnostic procedures complement the pri-	1107
1056	<i>ings of the International Conference on Learning</i>	mary DIF analysis.	1108
1057	<i>Representations (ICLR)</i> .	Local independence (Yen’s Q3). The MH proce-	1109
1058	Carol M. Woods. 2009. Empirical selection of anchors	dure assumes that item responses are conditionally	1110
1059	for tests of differential item functioning . <i>Applied</i>	independent given θ . Yen’s Q3 statistic measures	1111
1060	<i>Psychological Measurement</i> .	residual correlations between item pairs after con-	1112
1061	Siyuan Xu, Chao Zhang, and Wenjun Zhao. 2025.	ditioning on the latent trait. We compute Q3 for all	1113
1062	Benchmark self-evolving: A multi-agent framework	within-subject and cross-subject item pairs, with	1114
1063	for dynamic LLM benchmarking . In <i>Proceedings</i>	$ Q3 > 0.20$ as the threshold for local dependence.	1115
1064	<i>of the 63rd Annual Meeting of the Association for</i>	Elevated within-subject dependence would indicate	1116
1065	<i>Computational Linguistics (ACL)</i> .	that items within the same subject share variance	1117
1066	Jiahao Ying, Rui Cao, Mingsheng Luo, Shangqing	beyond what θ captures, which could inflate within-	1118
1067	Gao, Yuhang Zhang, Shuang Chen, and Jiawei Liu.	subject DIF rates.	1119
1068	2024. Automating dataset updates towards reliable	Domain-specificity. We test whether temporal	1120
1069	and timely evaluation of large language models. In	DIF concentrates in domains where models im-	1121
1070	<i>Findings of the Association for Computational Lin-</i>	proved most. At the ecological level, we com-	1122
1071	<i>guistics: ACL 2024</i> .	pute Spearman ρ between subject-level C% and	1123

subject-level accuracy improvement from 2023 to 2024. At the item level, logistic regression predicts item-level DIF C classification from domain-level accuracy change, with and without item property controls (difficulty, discrimination). This decomposition quantifies how much of the observed instability is attributable to aggregate domain-level gains versus item-specific behavioral change.

Instability decomposition. When domain accuracy improvement is included alongside item-level controls, we compute the counterfactual C% that would remain after removing domain-level effects. The gap between observed and counterfactual C% measures the contribution of domain-level improvement to the total instability.

A.2 Statistical Considerations

Multiple testing. The ETS C classification imposes a dual threshold (effect size $|\Delta_{\text{MH}}| \geq 1.5$ and significance $p < 0.05$), which provides implicit control over false positives. As a robustness check, we apply Benjamini-Hochberg FDR correction at $q = 0.05$. The random-split sanity check provides an empirical estimate of the false positive rate under the null hypothesis.

Threshold sensitivity. We report results across a range of DIF thresholds ($|\Delta_{\text{MH}}| \in \{1.0, 1.5, 2.0, 2.5\}$) to verify that qualitative conclusions do not depend on the specific ETS C cutoff. The standard threshold of $|\Delta_{\text{MH}}| \geq 1.5$ corresponds to the ETS C definition used throughout the paper.

IRT model fit. For transparency, we report that the MMLU response matrix shows poor fit to a unidimensional 2PL model (expected for a 57-subject multidimensional benchmark). This does not affect MH-DIF validity, which relies only on total score stratification rather than parametric IRT assumptions.

Architecture heterogeneity. The METABENCH MMLU population consists of 99.98% decoder-only models (5,226 of 5,227). This near-complete architectural homogeneity eliminates architecture as a confound for DIF analysis. We verify this by excluding the single encoder-decoder model and confirming identical results.

Power sensitivity to item discrimination. The Monte Carlo power analysis in Section 4.1 uses

average 2PL parameters ($\bar{a} = 1.2$, $\bar{b} = -0.84$). Table 4 reports TPR as a function of discrimination a (other parameters at their means, $N = 80$ per cohort, $|\Delta_b| = 1.4$, 1,000 replications). TPR increases monotonically from 0.40 ($a = 0.5$) to 0.95 ($a = 2.0$); FPR remains below 0.04 throughout. The MMLU item pool shows substantial heterogeneity: a ranges from 0.006 to 2.657 (SD = 0.54, IQR = $[0.75, 1.64]$). The aggregate TPR of 0.68 from Phase 2 reflects this full distribution, while TPR at mean $a \approx 1.2$ is 0.87—the gap is consistent with Jensen’s inequality applied to the concave $\text{TPR}(a)$ function. Although 2PL fit is poor for MMLU’s multidimensional structure, the power analysis serves as one leg of a triangulation: parametric estimates (TPR = 0.68–0.80) are validated by empirical semi-synthetic injection ($\Delta\text{TPR} = +0.508$ at $p = 0.5$; Section 4.5), confirming adequate detection power regardless of parametric assumptions.

Table 4: Power sensitivity to item discrimination. TPR = true positive rate for ETS C detection ($|\Delta_{\text{MH}}| \geq 1.5$) at $N = 80$ per cohort, $|\Delta_b| = 1.4$, 1,000 replications.

a	TPR	95% CI	FPR
0.5	0.398	[0.20, 0.60]	0.033
0.8	0.683	[0.45, 0.90]	0.036
1.0	0.808	[0.60, 0.95]	0.038
1.2 ($\approx \bar{a}$)	0.874	[0.70, 1.00]	0.038
1.5	0.930	[0.80, 1.00]	0.038
2.0	0.949	[0.80, 1.00]	0.037

Logistic regression specification. The Simpson’s paradox analysis (Section 4.2) uses binary logistic regression with DIF-C classification as the dependent variable, domain-level accuracy change (2024 minus 2023 mean accuracy) as the primary predictor, and item difficulty (proportion correct) and IRT discrimination (point-biserial correlation) as controls. The interaction between domain accuracy change and item difficulty is included. Non-uniform DIF detection (Section 4.4) uses logistic regression with a group $\times$ ability interaction term; a significant interaction ($p < 0.05$) indicates non-uniform DIF.

B Structural Analysis Details

Two structural controls and a sensitivity check fail to reduce the temporal C% below 31%, revealing a structural property of the data rather than a methodological limitation. Matched-ability quintile analysis restricts comparisons to models of

similar total score and yields enrichment ratios of 0.94–1.04, with temporal C% unchanged at approximately 0.31. Within-subject DIF analysis on 5 representative subjects, intended to isolate subject-specific effects, finds mean C% = 38.1%, exceeding the 25% threshold that would indicate viable bifactor structure. As a sensitivity check, IRT θ -conditioning with 20 ability strata achieves $r(\theta, \text{total score}) = 0.9986$ and C% = 0.312, confirming that the external total score already captures nearly all ability variance relevant to DIF computation. Table 2 summarizes these results.

The temporal FPR reflects genuine item-level instability that persists regardless of how ability is measured or conditioned. Yen’s Q3 diagnostic reveals mild local dependence within subjects (11.3% of within-subject item pairs exceed $|Q3| > 0.20$, vs. 2.2% cross-subject; $5\times$ ratio, $p = 2.3 \times 10^{-9}$), which contributes to elevated within-subject DIF rates but is not the primary driver: the cross-subject matching protocol addresses this bias, and temporal C% persists under external matching (full Q3 distributions available in the supplementary material).

Table 5: Domain-level Cohen’s d between 2023 and 2024 cohorts. The narrow range (0.12) indicates nearly uniform improvement across domains, supporting total-score matching adequacy for DIF analysis (Section 3).

Domain	Cohen’s d
STEM	0.261
Humanities	0.159
Social Science	0.147
Other	0.146

Table 6: Multi-subject injection robustness. k = number of simultaneously contaminated subjects (out of 57); values are mean $\pm$ SD across 5 random subject draws at exploitation rate $p = 0.5$.

k	ΔFPR	ΔTPR
1	0.000 $\pm$.000	0.371 $\pm$.048
3	0.003 $\pm$.002	0.385 $\pm$.031
5	0.007 $\pm$.003	0.392 $\pm$.027
10	0.047 $\pm$.022	0.401 $\pm$.024
20	0.069 $\pm$.018	0.410 $\pm$.020

C Design Effect Estimation

Model families (e.g., Mistral, Llama-3) induce intra-class correlation (ICC) that inflates Mantel-Haenszel χ^2 statistics beyond the nominal χ^2_1 reference distribution. We estimate the design effect

Table 7: Leave-one-family-out (LOFO) robustness on MMLU. Removing non-dominant families ($\leq 15\%$ of their cohort) changes C% by at most 2.93 pp. Larger shifts from removing Mistral or Llama-3 are driven almost entirely by family composition, not statistical power loss: Monte Carlo random-subsample calibration (200 replicates per scenario, matching the reduced sample size) shows the power effect is ≤ 1 pp, while the composition effect accounts for 95–113% of the observed C% drop (Mistral: -8.44 pp composition vs. -0.41 pp power; Llama-3: -8.17 pp vs. $+0.91$ pp; observed C% falls below the 95% random-subsample range in both cases). Stratification uses total test score for computational efficiency; cross-subject matching yields identical baseline C%.

Family	N removed	Cohort share	C% (LOFO)
<i>Full sample</i>	—	—	31.32
Mistral	547	51% ref	22.47
Llama-3	315	39% foc	24.06
Mixtral	120	15% foc	34.26
Llama-2	161	15% ref	31.48
Solar	91	11% foc	31.01

(DEFF) per item following Rao and Scott (1992): $\text{DEFF}_i = 1 + (\bar{m} - 1) \cdot \text{ICC}_i$, where $\bar{m}$ is the mean family size and ICC_i is the one-way random-effects ICC for item i across 30 families (restricted to ref+foc cohort; 1,884 models with family ≥ 2).

Two estimates of $\bar{m}$ bracket the adjustment: (1) *Simple mean* ($\bar{m} = 62.8$), treating all families equally, yields $\text{DEFF} = 3.87$ and adjusted C% = 29.2%; (2) *Weighted mean* ($\bar{m} = \sum n_f^2 / \sum n_f = 254.5$; Kish, 1965), reflecting the effective cluster size experienced by a randomly sampled model, yields $\text{DEFF} = 12.79$ and adjusted C% = 13.5%. Mean ICC = 0.047 (median 0.041, Q3 = 0.061, max = 0.219; 4.8% of items exceed 0.10). After dividing each item’s χ^2_{MH} by its item-specific DEFF and re-applying ETS thresholds, 13.5–29.2% of items retain Category C classification depending on the $\bar{m}$ variant—both far exceeding the 0.18% random-split baseline (75–162 $\times$).

D Uncertainty Quantification Comparison

Table 8: Comparison of uncertainty quantification methods for the 31.3% C% estimate on MMLU.

Method	Key assumption	Result
Model-level bootstrap	Models independent	95% CI [29.4%, 38.1%]
Family-level cluster bootstrap	Family is clustering unit	95% CI [26.7%, 53.2%]
DEFF (simple $\bar{m}$)	Equal-weight families	Adjusted C% = 29.2%
DEFF (weighted $\bar{m}$)	Size-proportional weighting	Adjusted C% = 13.5%
Graduated de-dup. ($K=20$)	Remove within-family duplicates	C% = 29.1%

The cluster bootstrap upper bound (53.2%) is driven by extreme bootstrap resamples that over-represent Mistral ($n = 547$) or Llama-3 ($n = 315$), producing atypical cohort compositions. Three independent methods—cluster bootstrap lower bound (26.7%), simple-mean DEFF adjustment (29.2%), and graduated de-duplication at $K=20$ (29.1%)—converge on ~ 27 – 29% , establishing a robust lower bound that remains $>150\times$ the random-split baseline (0.18%). The weighted-mean DEFF (13.5%) represents a conservative extreme that accounts for the effective cluster size experienced by a randomly sampled model; even this lower bound exceeds the baseline by $75\times$.

E Cross-Benchmark Comparison

Table 9: Cross-benchmark comparison of DIF analysis on MMLU and GSM8K. Both benchmarks exhibit temporal measurement instability with validated semi-synthetic detection power. Standard and external matching use different exploitation rates (p) because each benchmark’s optimal operating point differs; external matching results are compared at $p = 0.5$ for both. Subscript $\pm$ values denote standard deviations across 5 random injection seeds.

Metric	MMLU	GSM8K
Items	12,508	1,319
Models	5,227	6,702
Temporal C%	31.3%	18.9%
Sanity C%	0.18%	0.3%
ΔTPR ($p=0.7$, standard)	–	$+0.818 \pm .006$
ΔTPR ($p=0.5$, external)	$+0.508 \pm .027$	$+0.514 \pm .028$

F Additional Figures

G Item-Property Predictive Analysis

To test whether DIF susceptibility is a fixed item property, we trained logistic regression and LightGBM classifiers to predict ETS C classification from 17 leak-free features: ref-cohort difficulty and discrimination, 11 text features (question/option length, vocabulary metrics), subject metadata, and web-overlap labels (Li et al., 2024b). Features were computed exclusively from the reference (2023) cohort to prevent information leakage from the focal cohort. Five-fold subject-stratified cross-validation (no within-subject train/test overlap) yielded AUROC = 0.603 ± 0.020 (LR) and 0.599 ± 0.026 (LightGBM)—near chance. Web-overlap labels contributed negligibly (SHAP < 0.015), independently confirming the DIF–contamination orthog-




Figure 3: Cross-subject vs. within-subject matching for semi-synthetic DIF detection. Cross-subject matching eliminates FPR inflation ($\Delta\text{FPR} = 0.000$) while preserving detection power ($\Delta\text{TPR} > 0.3$ for 4/5 subjects at $p = 0.5$). Within-subject matching inflates FPR by $+0.186$ on average.

onality reported in Section 4.3. These results indicate that temporal DIF reflects emergent cohort–item interactions rather than fixed item defects, motivating post-hoc monitoring over pre-screening.




Figure 4: Downstream impact of DIF-based item filtering on model rankings. Despite $\rho \approx 0.997$ overall, individual model families show systematic rank shifts when evaluated on stable items only.




Figure 6: Ranking degradation under DIF-ordered vs. random item removal across six benchmarks. DIF-ordered removal (orange) consistently degrades ρ faster than random removal (blue band = mean $\pm$ 1 SD, 20 seeds). Vertical dashed lines mark the Monitor/Flag (15%) and Flag/Act (25%) thresholds from Section 5.2.




Figure 5: Exploratory: GSM1K contamination gap vs. DIF ranking advantage ($N = 33$; small sample, results should be interpreted cautiously). The negative correlation ($r = -0.478$, $p = 0.005$) indicates that models with larger contamination gaps benefit less from DIF-flagged items. Three of 33 models exceed the Cook's D threshold ($4/N = 0.12$); leave-one-out analysis confirms all individual-removal correlations remain significant ($r \in [-0.54, -0.37]$, all $p < 0.05$), and removing all three outliers yields $r = -0.38$ ($p = 0.036$, $N = 30$).