

Temporal Measurement Drift in LLM Benchmarks: A Large-Scale Psychometric Audit of 5,227 Models

Anonymous ACL submission

Abstract

LLM benchmark scores are routinely compared across model generations, yet the assumption that individual items measure the same construct over time is rarely tested. We apply Mantel-Haenszel Differential Item Functioning (DIF) analysis with ETS severity classification to six METABENCH benchmarks (5,227 models, 27,125 total items), and find that 17.9–31.3% of items exhibit large DIF (ETS Category C; MMLU 95% bootstrap CI: [29.4%, 38.1%]) between 2023 and 2024 model cohorts—knowledge/reasoning benchmarks show higher DIF rates ($C\% \geq 26.9\%$) than procedural ones ($C\% \leq 18.9\%$), each 45–174 $\times$ above random baselines. This does not threaten benchmark validity ($\rho \approx 0.997$ between full and stable-item rankings) but reveals which capability dimensions shift across generations. Item-level decomposition shows 92.4% of the DIF signal is item-specific; a Simpson’s paradox in odds ratio regression confirms item properties, not accuracy gains, drive DIF. We expose a validation methodology gap: web-overlap labels measure global exposure, not differential functioning (enrichment ≈ 0.96). A semi-synthetic injection framework with cross-subject matching preserves detection power ($\Delta\text{TPR} = +0.508$ at $p = 0.5$) with zero false positive inflation. DIF-flagged items concentrate cross-generational discriminative information ($z = -8\sigma$ to -49σ vs. random removal), and a difficulty–direction interaction ($\rho = -0.597$) reveals that capability gains concentrate on hard items while easy items erode, suggesting DIF captures temporal construct shift rather than noise. LLM benchmarks need psychometric monitoring, not just score reporting.

1 Introduction

When we say “Model B outperforms Model A on MMLU,” we implicitly assume that the benchmark measures the same construct in the same way for

both models. This assumption of temporal comparability underpins every cross-generation comparison published on leaderboards and in model reports, yet it is rarely tested. We apply Mantel-Haenszel Differential Item Functioning (MH-DIF) analysis (Holland and Thayer, 1988) with Educational Testing Service (ETS) severity classification to the METABENCH response matrix (Kipnis et al., 2024),¹ and find that 31% of MMLU (Hendrycks et al., 2021) items receive ETS Category C classification ($|\Delta_{\text{MH}}| \geq 1.5$, $p < 0.05$), indicating large shifts in item functioning between models released in 2023 and 2024. This is not statistical noise: a random-split control within the 2023 cohort yields only 0.18% ETS C items. This pattern persists across all 57 MMLU subjects ($C\%$ range 21–48%, unweighted mean $\approx 30\%$), and item-level regression shows that 92.4% of the DIF signal is unrelated to aggregate domain-level accuracy gains. Nearly one in three MMLU items behaves differently depending on when the evaluated models were trained. Crucially, this measurement drift does not threaten benchmark validity: Spearman rank correlation between model rankings computed on all items versus stable-only items is $\rho \approx 0.997$. The drift is diagnostic, not destructive; it reveals which capability dimensions shift across model generations rather than rendering MMLU scores unreliable.

A growing body of contamination detection research has produced sophisticated item-level methods, from membership inference approaches such as Min-K% (Shi et al., 2024) to performance-based detectors such as ConStat (Dekoninck et al., 2024), DVD (Liang, 2026), and Self-Critique (Tao et al., 2026). These methods validate detection accuracy primarily against web-overlap labels such as those of Li et al. (2024b), which classify bench-

¹Analysis code and DIF audit scripts are available at <https://anonymous.4open.science/r/dif-audit>.




Figure 1: Global vs. differential contamination, illustrated via Item Characteristic Curves (ICCs). **Left:** global contamination produces no DIF; both cohorts encounter the item, and their ICCs overlap. **Right:** differential contamination produces DIF; the focal cohort’s ICC shifts because only that cohort was disproportionately exposed during training. The enrichment ratio of ≈ 1.0 between DIF flags and web-overlap labels is expected because most MMLU items fall into the left category.

mark items as contaminated based on verbatim or near-verbatim matches in publicly accessible training corpora. Such labels measure *global exposure*: whether an item exists in the data ecosystem that all training pipelines can access. They do not measure *differential item functioning*: whether an item behaves differently across model cohorts that may have had varying degrees of exposure or different training methodologies. Figure 1 illustrates the distinction. Most MMLU items are globally contaminated in the sense that they exist in web-accessible text, so all model cohorts encounter them and no differential signal arises. The items that function differentially are those where exposure or training methodology changed asymmetrically across cohorts. Validating differential detectors against global-exposure labels therefore conflates two phenomena. Direct comparison confirms near-independence between the two flag sets (enrichment = 0.964, Jaccard = 0.206, ϕ = 0.012), consistent across contamination subtypes.

We apply classical MH-DIF with ETS classification to the full METABENCH response matrix (5,227 models $\times$ 12,508 MMLU items), to our knowledge, the largest-scale DIF analysis on LLM benchmarks. This psychometric approach complements concurrent work on anchor-item calibration for score comparability (Habba et al., 2026) and addresses the fragility of existing contamination detectors under post-training perturbation (Wang et al., 2026). Semi-synthetic injection experiments confirm that MH-DIF is sensitive to differential

item behavior when present ($\Delta\text{TPR} = +0.508$ at exploitation rate $p = 0.5$), and our cross-subject matching protocol eliminates false positive inflation ($\Delta\text{FPR} = 0.000$). Three structural controls (matched-ability quintile analysis, IRT θ -conditioning, and within-subject pre-checks) confirm that the approximately 31% temporal DIF rate cannot be attributed to methodological artifacts, establishing it as a robust empirical finding about LLM benchmark measurement properties, replicated across all six METABENCH benchmarks (Section 4.7). MH-DIF detects *that* items function differently across cohorts, not *why*—the signal may reflect contamination, capability change, or training methodology shifts.

MMLU’s 12,508 items provide sufficient measurement redundancy that removing nearly a third does not alter rank order ($\rho \approx 0.997$). DIF should therefore be understood as a diagnostic lens rather than a quality filter: DIF-flagged items concentrate discriminative signal ($z = -8\sigma$ to -49σ vs. random removal), and model families show distinctive DIF profiles (Llama-3 gains +53 median rank positions, Mistral loses 35).

Our contributions are empirical and methodological, not algorithmic:

- Temporal measurement drift at scale:** Approximately 31% of MMLU items exhibit differential functioning across model generations under ETS C classification; replication across all six METABENCH benchmarks yields $C\% = 17.9\text{--}31.3\%$ ($45\text{--}174\times$ above random baselines), with knowledge/reasoning benchmarks showing higher DIF rates than procedural ones. DIF-flagged items concentrate discriminative signal ($z = -8\sigma$ to -49σ vs. random removal) and exhibit a difficulty-direction interaction ($\rho = -0.597$): gains on hard items, erosion on easy ones (Sections 4.2, 4.6 and 4.7).
- Ground truth meta-critique:** Web-overlap contamination labels (Li et al., 2024b) do not provide appropriate ground truth for item-level differential detectors. Direct empirical comparison on 9,229 items shows near-independence (enrichment = 0.964, Jaccard = 0.206, ϕ = 0.012), exposing a validation methodology gap (Section 4.3).
- Semi-synthetic validation framework:** Controlled injection with cross-subject matching

achieves $\Delta\text{FPR} = 0.000$ while preserving detection power ($\Delta\text{TPR} = +0.508$ at $p = 0.5$); cross-benchmark replication on GSM8K confirms generality (Sections 4.5 and 4.7).

2 Related Work

Contamination Detection. A growing family of methods detects benchmark contamination through model behavior or training data analysis. Membership inference approaches estimate whether a model has seen specific items during training: Min-K% (Shi et al., 2024) and Min-K%++ (Zhang et al., 2024) use token-level log-likelihood statistics, while Time Travel (Golchin and Surdeanu, 2024) probes models with temporal cues. Performance-based detectors identify contamination through statistical anomalies: ConStat (Dekoning et al., 2024) tests performance gaps between original and perturbed items, DVD (Liang, 2026) uses distributional verification, and Self-Critique (Tao et al., 2026) employs model self-evaluation. Balloccu et al. (2024) provide a broader taxonomy of leakage detection strategies. These methods share a common validation practice: they verify detection accuracy against web-overlap or n-gram labels (Li et al., 2024b) that classify items based on verbatim matches in publicly accessible corpora. Such labels measure *global exposure* (whether an item exists in training data) rather than *differential item functioning* (whether an item behaves differently across model cohorts). This global-vs-differential distinction motivates our ground truth analysis (Section 4.3). Work on LLM memorization (Carlini et al., 2023; Tirumala et al., 2022) establishes that memorization depends on data frequency and model scale, reinforcing the distinction between global exposure and differential behavior. Wang et al. (2026) show that existing contamination detectors are fragile under reinforcement learning, further motivating behavioral approaches grounded in psychometric theory.

IRT in LLM Evaluation. IRT adoption in LLM evaluation has accelerated: Lalor et al. (2019) introduced IRT to NLP, Vania et al. (2021) assessed test set discriminability, and Rodriguez et al. (2021) demonstrated that not all evaluation examples are equally informative. Recent extensions include PSN-IRT (Zhou et al., 2026) for benchmark quality at scale, TinyBenchmarks (Maia Polo et al., 2024) for IRT-based item selection, scalability coefficients for detecting bad benchmark items (Hardy

et al., 2026), ATLAS (Li et al., 2025) for computerized adaptive testing of LLMs, and competency-focused IRT evaluation in medical domains (Luo et al., 2025). The closest concurrent work is Growing Pains (Habba et al., 2026), which addresses score comparability through fixed-parameter calibration (FPC; Kim and Kolen, 2019). Our work is complementary: Growing Pains assumes anchor item stability, while DIF analysis diagnoses which items violate this assumption.

DIF Methodology and Measurement Invariance.

The Mantel-Haenszel procedure for DIF detection was introduced by Holland and Thayer (1988) and remains the most widely used method in educational testing due to its simplicity and well-characterized statistical properties (Clauser and Mazor, 1998). Alternative approaches include logistic regression DIF (Swaminathan and Rogers, 1990), SIBTEST (Shealy and Stout, 1993), Lord’s χ^2 test (Lord, 1980), and the Raju area method (Raju, 1988); Zumbo (1999) provides a comprehensive comparison. DIF analysis is a specific instance of measurement invariance testing (Meredith, 1993; Vandenberg and Lance, 2000), which examines whether a measurement instrument functions equivalently across populations. MIMIC models offer a latent-variable approach to the same problem (Muthén, 1985; Woods, 2009), and Teresi and Fleishman (2007) reviews DIF methods across applied domains. Cross-subject matching—computing the stratification variable from domains other than the one being tested—has psychometric precedent in the DIF literature (Nandakumar, 1993; Donoghue and Allen, 1993); our adaptation addresses LLM-specific challenges of multi-domain local dependence and simultaneous contamination control at scale. We do not modify the MH procedure; our contribution is the application context and the empirical findings it reveals. The scale of this analysis (5,227 examinees across a 57-subject multidimensional ability space) exceeds typical educational testing applications, where Belzak (2020) note that DIF methods can behave differently at extreme sample sizes. The primary methodological challenge is the impact-versus-DIF confound: genuine ability differences between temporal cohorts can inflate DIF estimates. We address this through matched-ability quintile analysis, IRT θ -conditioning, and a random-split control (Section 4.2).

Benchmark Quality and Design. The Benchmark Health Index (Zhu et al., 2026) and Benchmark² (Qian and Huang, 2026) assess static benchmark quality; our temporal DIF analysis is complementary, diagnosing *when* measurement properties shift rather than evaluating them at a single snapshot. Maeda (2025) predict which item features associate with DIF in educational settings; we instead detect DIF empirically across LLM cohorts at scale. Dynamic benchmarks (Kiela et al., 2021; Li et al., 2024a; White, 2025) mitigate contamination by construction, yet none provides psychometric equating across versions. Chatbot Arena (Chiang et al., 2024) sidesteps static benchmarks through pairwise preference. Broader surveys (Chang et al., 2024; Burnell et al., 2023) identify benchmark saturation and lack of item-level transparency as systemic challenges that our DIF analysis directly addresses.

3 Method

We adopt Differential Item Functioning (DIF) analysis from educational measurement to assess whether individual benchmark items behave consistently across groups of LLMs. The procedure requires only a binary response matrix and group labels, with no access to model internals or training data. Below we define the data, the statistical test, the cohort designs, and a semi-synthetic validation framework.

3.1 Data

The primary dataset is the METABENCH response matrix (Kipnis et al., 2024), which records binary outcomes (1 = correct, 0 = incorrect) for 5,227 models evaluated on 12,508 MMLU items (Hendrycks et al., 2021). Models span release dates from 2023 to 2024, drawn from the Open LLM Leaderboard. MMLU covers 57 subjects grouped into five broad domains (STEM, Humanities, Social Sciences, Other, and a mixed category). We supplement the response matrix with the web-overlap contamination labels of Li et al. (2024b), which classify 9,229 of 12,508 items (73.8% coverage) as “contaminated” or “clean” based on verbatim matches in publicly accessible web corpora. These labels serve as an external reference for evaluating the relationship between DIF flags and existing contamination indicators.

3.2 Mantel-Haenszel DIF and ETS Classification

In educational testing, Differential Item Functioning (DIF) occurs when an item’s statistical relationship to the latent trait θ differs between two groups after conditioning on ability level. An item with DIF is not merely harder for one group; it functions differently relative to what the group’s overall ability would predict. The standard terminology refers to test-takers as “examinees” and questions as “items.” In our setting, examinees are LLMs and items are benchmark questions.

The Mantel-Haenszel (MH) procedure (Holland and Thayer, 1988) detects DIF by stratifying examinees into K score bins based on total score (a proxy for θ) and computing a common odds ratio across strata. We use $K = 5$ equiprobable strata (quintile-based binning) following ETS recommended practice for large-scale assessments (Donoghue and Allen, 1993; Clauser and Mazor, 1998). Within each stratum k , the responses of the reference group and the focal group form a 2×2 contingency table:

$$\alpha_{\text{MH}} = \frac{\sum_{k=1}^K A_k D_k / T_k}{\sum_{k=1}^K B_k C_k / T_k}, \quad (1)$$

where A_k and B_k are the number of correct and incorrect responses in the reference group for stratum k , C_k and D_k are the corresponding counts for the focal group, and T_k is the total number of examinees in stratum k . Under the null hypothesis of no DIF, $\alpha_{\text{MH}} = 1$. The log odds ratio is converted to the ETS delta scale:

$$\Delta_{\text{MH}} = -2.35 \ln(\alpha_{\text{MH}}). \quad (2)$$

The Educational Testing Service (ETS) classifies DIF severity into three categories based on $|\Delta_{\text{MH}}|$. Category A (negligible) requires $|\Delta_{\text{MH}}| < 1.0$. Category C (large) requires both $|\Delta_{\text{MH}}| \geq 1.5$ and statistical significance at $p < 0.05$ from the MH chi-squared test. Category B (moderate) includes all remaining items: those with $1.0 \leq |\Delta_{\text{MH}}| < 1.5$, or $|\Delta_{\text{MH}}| \geq 1.5$ without statistical significance. The dual threshold in ETS C (effect size *and* significance) guards against flagging items that show large but unreliable effect sizes. Throughout this paper, we report the percentage of items classified as ETS C, denoted C%. For example, if

2024 models consistently outperform 2023 models on a formal-logic item *after controlling for overall ability*, the pooled odds ratio $\alpha_{\text{MH}} = 0.42$ yields $\Delta_{\text{MH}} = -2.35 \ln(0.42) = 2.04$, classified as ETS C ($|\Delta_{\text{MH}}| \geq 1.5$, $p < 0.001$).

The choice of matching variable (the total score used for stratification) affects both validity and interpretation. We distinguish three matching strategies: (1) *standard matching* uses the within-benchmark total score on all items of the same benchmark; (2) *cross-subject matching* uses the total score on all subjects *other than* the one containing the item under test, preventing the matching variable from absorbing injected contamination signals (Section 3.4); and (3) *external matching* uses the total score on a different benchmark entirely (e.g., MMLU total score as the matching variable when testing DIF on GSM8K items), providing an independent ability proxy uncontaminated by the target benchmark. Cross-subject and external matching both address matching variable contamination but at different granularities: cross-subject operates within a multi-domain benchmark, while external operates across benchmarks.

LLM-specific challenges. Applying MH-DIF to LLMs introduces assumptions absent in educational testing. First, models trained on overlapping corpora violate local independence; Yen’s Q3 analysis confirms this is bounded ($r = 0.09$, $p = 0.50$; Section A). Second, fine-tuned variants and merges within model families are not independent examinees; graduated de-duplication shows family structure does not inflate C% (Section 4.4). Third, the ability distribution is non-normal (bimodal from base vs. instruction-tuned models), motivating quintile-based rather than parametric stratification. Fourth, MMLU’s 57 subjects span multiple latent dimensions, while MH assumes unidimensionality; cross-subject matching and within-subject DIF checks address this (Section 4.4).

3.3 Cohort Designs

We construct five cohort comparisons: (1) *Temporal*: 2023 ($N = 1,080$) vs. 2024 ($N = 807$) models, the primary comparison; (2) *Ability-quintile*: Q1–Q2 vs. Q4–Q5 within the same temporal window, testing whether ability differences alone produce DIF; (3) *Within-family*: Llama-2 ($N = 161$) vs. Llama-3 ($N = 308$), isolating generational changes within a single family; (4) *Within-subject*: per-subject DIF to test domain concentration; and

(5) *Random split*: two random halves of the 2023 cohort ($N = 540$ each) as a sanity check.

3.4 Semi-Synthetic Injection Framework

To evaluate whether MH-DIF can detect differential contamination when it is present, we construct a semi-synthetic framework with controlled ground truth.

Procedure. From items classified as ETS A in the uninjected data (negligible DIF), we select 20% uniformly at random as the “contaminated” set. For each contaminated item, we flip incorrect responses to correct in the focal group at exploitation rate $p \in \{0.0, 0.3, 0.5, 0.7, 1.0\}$. We then run MH-DIF on the modified response matrix and measure detection performance. We report $\Delta\text{TPR} = \text{TPR}(\text{injected}) - \text{TPR}(\text{baseline})$ and $\Delta\text{FPR} = \text{FPR}(\text{injected}) - \text{FPR}(\text{baseline})$, where TPR is the true positive rate on injected items and FPR is the false positive rate on uninjected items. Each injection level is repeated across 5 random seeds for uncertainty estimation.

Cross-subject matching. The choice of matching variable is critical for valid semi-synthetic evaluation. When contamination is injected into items within subject X , the total score on subject X absorbs part of the injection signal. Using this contaminated total score as the MH matching variable violates the conditional independence assumption and inflates FPR. Cross-subject matching resolves this by computing the matching variable from all subjects *other than* the one being tested.

Proposition 1 (Zero FPR under single-subject cross-subject matching). *Let $M_i^{(-s)} = \sum_{j \neq s} Y_{ij}$ be the matching variable for examinee i on focal subject s . Under single-subject injection (contamination confined to subject s), the tested item belongs to s and does not contribute to $M_i^{(-s)}$; injecting contamination into s leaves $M_i^{(-s)}$ unchanged. Therefore $\Delta\text{FPR} = 0$ by construction.*

When contamination spans multiple subjects simultaneously, $M_i^{(-s)}$ may absorb signal from other contaminated subjects, potentially inflating FPR. We evaluate this multi-subject scenario empirically in Section A.

Empirically, $\Delta\text{FPR} = 0.000$ across all five MMLU domains (bootstrap 95% CI = $[0.000, 0.000]$, 5,000 resamples), while detection power is preserved ($\Delta\text{TPR} > 0.3$ at $p = 0.5$ for four of five subjects). We adopt cross-subject

matching as the standard protocol for all semi-synthetic analyses.

Diagnostic procedures (Yen’s Q3 local independence, domain-specificity analysis, DIF decomposition) and statistical details (multiple testing correction, threshold sensitivity, architecture heterogeneity checks) appear in Section A.

4 Experiments

We apply MH-DIF to the METABENCH response matrix (Kipnis et al., 2024) (5,227 models $\times$ 12,508 MMLU items (Hendrycks et al., 2021)), with temporal cohorts defined by release date (2023 vs. 2024). Cross-benchmark validation across all six METABENCH benchmarks appears in Section 4.7.

4.1 Validation: DIF Detection Power

A Monte Carlo power analysis (1,000 replications per condition) using empirical 2PL parameters from PSN-IRT (Zhou et al., 2026) ($\bar{a} = 1.2$, $\bar{b} = -0.84$; 14,042 MMLU items) confirms that at $N \geq 80$ per cohort, ETS C achieves TPR of 0.68–0.80, FPR of 0.036–0.042, and AUC of 0.93–0.94 for items with $|\Delta_{\text{MH}}| \geq 1.4$ (details in Section A). Our temporal cohorts ($N_{\text{ref}} = 1,080$, $N_{\text{focal}} = 807$) far exceed this threshold.

4.2 Temporal Comparability: 31% Measurement Drift

We compare 2023 ($N = 1,080$) against 2024 ($N = 807$) models; 31.3% of items receive ETS C classification ($|\Delta_{\text{MH}}| \geq 1.5$; 95% bootstrap CI: [29.4%, 38.1%], $B = 1,000$ model-level resamples). Benjamini-Hochberg FDR correction at $q = 0.05$ removes only 3 items (31.32% $\rightarrow$ 31.30%), because the dual ETS threshold on effect size and significance already provides implicit multiple-testing protection.

Three controls confirm this reflects genuine temporal change. A random split of the 2023 cohort yields only 0.18% ETS C items (sampling noise), and an ability-quintile comparison within the same temporal cohort yields 2.9% (ability differences alone). A within-family comparison (Llama-2 vs. Llama-3) produces 61.3% ETS C items, elevated by large ability gaps that confound the DIF signal (Holland and Thayer, 1988). The contrast between 2.9% (ability-controlled) and 61.3% (ability-confounded) brackets the temporal rate of 31.3%. Crucially, the temporal cohort ability gap is negligible: Cohen’s $d = -0.17$ (below the $|d| = 0.2$

Table 1: DIF landscape across cohort definitions. ETS C denotes the percentage of items with $|\Delta_{\text{MH}}| \geq 1.5$. FPR is measured against Li et al. web-overlap labels on the “clean” (non-overlapping) subset. Enrichment is the ratio of observed to expected overlap between DIF flags and web-overlap labels.

Cohort Definition	$N_{\text{ref}} / N_{\text{focal}}$	C%	FPR (clean)	Enrich.
Temporal (2023 vs. 2024)	1,080 / 807	31.3	0.318	0.964
Ability quintile (Q1–2 vs. Q4–5)	$\sim 2,600 / \sim 2,600$	2.9	0.027	1.19
Within-family (Llama-2 vs. 3)	161 / 308	61.3	–	–
Random split (sanity)	540 / 540	0.18	–	–
Within-subject (avg. 5 subj.)	<i>varies</i>	38.1	–	–

small-effect threshold) with 76.4% distributional overlap. Under strict 1:1 nearest-neighbor matching ($d \approx -0.010$, $N = 638$ per group, 5 seeds), $C\% = 28.9\% \pm 0.1\%$ —a 7.7% relative reduction from 31.3%, confirming that at most ~ 2.4 pp of observed DIF is attributable to ability mismatch.

Across all 57 subjects, $C\%$ ranges from 20.7% to 48.4% (mean $\approx 30\%$; Table 1). DIF rates are domain-invariant (Kruskal-Wallis $p = 0.78$), but DIF *direction* exhibits domain dependence: STEM items favor 2024 models (55.4%), Humanities favor 2023 (64.4%). Item-level regression shows 92.4% of the DIF signal is item-specific (pseudo $R^2 = 0.0007$). A Simpson’s paradox confirms item properties—not accuracy gains—drive DIF: domain accuracy positively predicts DIF at the aggregate level ($\text{OR} = 1.55$, $p = 0.001$) but reverses after controlling for item difficulty and discrimination ($\text{OR} = 0.73$, $p = 0.030$). The conclusion is robust across thresholds: $C\% = 48.9\%$ at $|\Delta_{\text{MH}}| \geq 1.0$, 18.3% at ≥ 2.0 , 10.2% at ≥ 2.5 ; ETS C’s dual threshold makes results less sensitive to the number of strata K than significance-only approaches ($K = 5$ per ETS practice (Donoghue and Allen, 1993)).

4.3 Ground Truth Mismatch

Direct comparison of DIF flags against the web-overlap labels of Li et al. (2024b) on 9,229 annotated items confirms that the two methods capture orthogonal signals. The enrichment ratio is 0.964 (Table 1): DIF-C rates are 30.7% among web-contaminated items and 31.8% among clean items. Association tests yield $\phi = 0.012$, $\chi^2 = 1.28$ ($p = 0.258$), and $\text{OR} = 1.05$ [0.96, 1.15]; the Jaccard index between flagged sets is 0.206. This near-independence holds across contamination subtypes (input-only enrichment = 0.986, input-and-label = 0.960). The null result does not indicate detection failure: semi-synthetic injection produces a clear dose-response ($\Delta \text{TPR} = +0.508$ at $p = 0.5$;




Figure 2: Semi-synthetic dose-response for DIF detection on MMLU. ΔTPR increases monotonically with exploitation rate p , confirming that MH-DIF is sensitive to differential item behavior when present. Baseline (no injection) produces $\Delta\text{TPR} \approx 0$.

Figure 2). The disconnect arises because web-overlap labels measure *global exposure* (whether an item exists in training corpora), while DIF measures *differential functioning* (whether an item’s relationship to ability changes across cohorts). The Li et al. labels are the wrong ground truth for a differential method. This near-independence extends to behavioral contamination detectors: the enrichment ratio of 0.964 and Jaccard overlap of 0.206 between DIF flags and web-overlap labels imply that any detector validated against these labels—including ConStat (Dekoninck et al., 2024)—is evaluated on a ground truth orthogonal to differential functioning.

4.4 Structural Analysis

Three structural controls fail to reduce temporal C% below 31%: matched-ability quintiles, IRT θ -conditioning ($r(\theta, \text{total score}) = 0.9986$), and within-subject DIF (C% = 38.1%). Logistic regression DIF confirms 92.1% concordance with MH and reveals additional non-uniform DIF (24.8%), making 31% conservative. Graduated deduplication ($K=10$: 22.7%; $K=20$: 29.1%) confirms family structure does not inflate C%. Ability-matched analysis reduces Llama-2 vs. Llama-3 C% from 61.3% to 28.3% and temporal C% from

31.3% to 28.9%—consistent matched values (~ 28 –29%) despite vastly different unmatched rates, suggesting genuine item-level DIF accounts for ~ 28 –29% of items. Six alternative explanations (ability confounding, family inflation, θ -conditioning, within-subject effects, ranking impact, local dependence) were tested; none reduces C% below $\sim 28\%$ (Section B).

4.5 Semi-Synthetic Framework Validation

The critical methodological finding is that the choice of matching strategy determines validity. Cross-subject matching computes the conditioning variable from subjects *other than* the one being tested, achieving $\Delta\text{FPR} = 0.000$ across all five MMLU subjects while four of five produce $\Delta\text{TPR} > 0.3$ at $p = 0.5$. Within-subject matching inflates FPR by +0.186 on average because the matching variable absorbs part of the injection signal. We adopt cross-subject matching as the standard protocol (Figure 3).

4.6 Downstream Impact

Spearman correlation between full and stable-item rankings is $\rho \approx 0.997$: aggregate ordering is preserved despite 31% measurement drift. DIF-flagged items nonetheless concentrate discriminative signal ($z = -8\sigma$ to -49σ vs. random removal), though a circularity ablation confirms DIF-based removal is the *least* disruptive non-random strategy. Model families show distinctive DIF profiles: Llama-3 derivatives gain +53 median rank positions, Mistral loses 35 (Figure 4). Removing C-flagged items disproportionately disadvantages 2024 models (cohort gap = 96.2 positions, $40\times$ larger than random removal; $p < 0.001$), confirming that DIF items carry concentrated inter-generational signal.

4.7 Cross-Benchmark Validation

Temporal DIF extends across all six METABENCH benchmarks (Table 2). C% ranges from 17.9% (HellaSwag) to 31.3% (MMLU), each 45 – $174\times$ above its random-split baseline—none attributable to sampling noise. A task-type dichotomy emerges: knowledge and reasoning benchmarks (MMLU, WinoGrande (Sakaguchi et al., 2020), TruthfulQA (Lin et al., 2022), ARC (Clark et al., 2018)) show C% = 26.9–31.3%, while procedural benchmarks (GSM8K, HellaSwag (Zellers et al., 2019)) show C% = 17.9–18.9%, consistent with knowledge-intensive items being more sensitive to

Table 2: Temporal DIF across six METABENCH benchmarks. C% = ETS C items ($|\Delta_{MH}| \geq 1.5$) between 2023 and 2024 cohorts. All six show substantial DIF (45–174 $\times$ above random baseline), with knowledge/reasoning benchmarks exhibiting higher C% than procedural ones.

Benchmark	Items	C%	Rand. C%	Ratio
MMLU	12,508	31.3	0.18	174 $\times$
WinoGrande	1,267	30.3	0.32	95 $\times$
TruthfulQA	817	29.7	0.37	80 $\times$
ARC	1,172	26.9	0.60	45 $\times$
GSM8K	1,319	18.9	0.30	63 $\times$
HellaSwag	10,042	17.9	0.28	64 $\times$

shifts in training data composition across model generations.

On GSM8K specifically, semi-synthetic injection with external matching confirms detection power ($\Delta\text{TPR} = +0.514$ at $p = 0.5$; Section D). An exploratory GSM1K analysis ($N = 33$) shows a negative correlation between contamination gap and DIF ranking advantage ($r = -0.478$, $p = 0.005$; partial $r = -0.385$, $p = 0.027$), suggesting DIF captures capability change distinct from memorization (Figure 5).

Even within a 6-month window (2023-H2 vs. 2024-H1), DIF remains substantial (MMLU C% = 34.35%, 181 $\times$; TruthfulQA C% = 31.33%, 51 $\times$ above random baselines).

5 Discussion

5.1 Implications for Benchmark Temporal Comparability

The 31% ETS C DIF rate extends across all six METABENCH benchmarks (C% = 17.9–31.3%, 45–174 $\times$ above baselines; Table 2), with knowledge/reasoning benchmarks more affected (C% $\geq 26.9\%$) than procedural ones (C% $\leq 18.9\%$). The Simpson’s paradox (Section 4.2) confirms item properties, not accuracy gains, drive DIF. This does not invalidate ranking ($\rho \approx 0.997$) but signals measurement drift across model generations. A fundamental limitation is that MH-DIF cannot distinguish contamination-driven DIF from legitimate capability evolution: the observed 31% rate represents an upper bound on measurement degradation and a lower bound on genuine construct change. Disentangling the two requires access to training data or controlled ablations rarely available for open-weight models.

5.2 Ground Truth Meta-Critique

The near-independence between DIF flags and web-overlap labels (enrichment = 0.964, Jaccard = 0.206, $\phi = 0.012$) admits two interpretations: web-overlap labels may be the wrong ground truth for a differential method, or DIF may fail to detect contamination. The semi-synthetic dose-response ($\Delta\text{TPR} = +0.508$ at $p = 0.5$) partially resolves this by confirming that MH-DIF *can* detect differential item behavior when present; the near-independence therefore reflects a category error in the ground truth rather than a detection failure.

5.3 DIF as Diagnostic Lens

The GSM1K correlation ($r = -0.478$, partial $r = -0.385$; Section 4.7) provides exploratory evidence that DIF captures behavioral change distinct from memorization. DIF-C items are harder ($d = -0.247$) and more discriminating ($d = 0.230$), with a difficulty–direction interaction ($\rho = -0.597$): gains on hard items, erosion on easy ones. Yen’s Q3 confirms local dependence does not drive DIF ($r = 0.09$, $p = 0.50$).

5.4 Recommended DIF Audit Protocol

These findings motivate a benchmark maintenance protocol triggered whenever ≥ 500 new models accumulate, using cross-subject matching. We propose graduated thresholds: C% < 15% $\rightarrow$ monitor (log and re-audit next cycle); 15–25% $\rightarrow$ flag (label consecutively flagged items “drift-prone,” flip-flopping items “volatile”); > 25% $\rightarrow$ act (item review, rotate or retire persistently flagged items). MMLU’s C% = 31.3% exceeds the action threshold, indicating that item-level review is warranted.

6 Conclusion

MH-DIF analysis of six METABENCH benchmarks (27,125 items, 5,227 models) reveals 17.9–31.3% temporal measurement drift (ETS C); a Simpson’s paradox confirms item properties, not accuracy gains, drive DIF. Web-overlap labels are near-independent of DIF flags (enrichment ≈ 0.96); our semi-synthetic framework validates detection power ($\Delta\text{FPR} = 0.000$). DIF-flagged items concentrate discriminative information ($z = -8\sigma$ to -49σ) with a difficulty–direction interaction ($\rho = -0.597$): gains on hard items, erosion on easy ones. LLM benchmarks need psychometric monitoring alongside score reporting.

Limitations

Methodological scope of MH-DIF. The Mantel-Haenszel procedure detects only *uniform* DIF, where the odds ratio between cohorts is constant across ability strata. Logistic regression DIF analysis reveals an additional 24.8% of items exhibiting non-uniform DIF ($|\beta_{\text{interaction}}| \geq 0.5$), making the 31% ETS C estimate conservative. Future work should apply non-uniform DIF methods (e.g., logistic regression or IRT-based approaches) as the primary analysis to capture the full spectrum of measurement drift. Critically, MH-DIF is nonparametric and conditions on observed total score, not latent θ —its validity does not depend on IRT model fit (Holland and Thayer, 1988). Our Yen’s Q3 analysis confirms that local dependence domains do not correlate with DIF domains ($r = 0.09$, $p = 0.50$), further supporting the robustness of the MH procedure in this setting.

Ability confounding. Although the temporal cohort ability gap is negligible (Cohen’s $d = -0.17$, overlap = 76.4%) and nearest-neighbor matching confirms that at most 2.4 percentage points of C% are attributable to ability mismatch, within-family comparisons with large ability gaps (e.g., Llama-2 vs. Llama-3, $d = -0.99$) produce substantially inflated C% (61.3% unmatched vs. 28.3% matched). Any application of MH-DIF to cohorts with larger ability differences must incorporate ability matching or stratification beyond the standard 5-strata procedure.

Sample size requirements. Our power analysis indicates reliable ETS C detection requires $N \geq 80$ per cohort. The temporal cohorts (1,080 and 807 models) far exceed this threshold, but the GSM1K correlation analysis ($N = 33$) has limited statistical power (95% CI = $[-0.71, -0.16]$). The negative correlation between contamination gap and DIF ranking advantage ($r = -0.478$) is exploratory and warrants replication with larger model samples.

Contamination vs. capability change. DIF analysis detects *that* items function differently across cohorts but cannot determine *why*. The differential signal may reflect contamination, capability change, training methodology shifts, or interactions among these factors. We present evidence that DIF captures something distinct from memorization (negative GSM1K correlation, semi-synthetic

dose-response), but a definitive causal decomposition remains out of reach without access to model training data.

Benchmark scope. We analyze six METABENCH benchmarks (MMLU, WinoGrande, TruthfulQA, ARC, GSM8K, HellaSwag; 27,125 total items), all multiple-choice or short-answer benchmarks scored as binary correct/incorrect. The generalizability to open-ended generation benchmarks (e.g., AlpacaEval, MT-Bench), coding benchmarks (e.g., HumanEval), or benchmarks with non-binary scoring remains untested. The binary response format is a prerequisite for MH-DIF; extending to continuous or partial-credit scoring would require polytomous DIF methods.

Temporal cohort definition. We define cohorts by calendar year of model release (2023 vs. 2024), a coarse grouping that aggregates models with diverse training pipelines, data mixtures, and post-training procedures. Finer-grained cohort definitions (e.g., by training data cutoff, model family, or alignment method) might reveal more targeted patterns of differential functioning but would reduce per-cohort sample sizes and statistical power.

References

- Simone Balloccu, Patrícia Schmidtova, Mateusz Lango, and Ondrej Dusek. 2024. [Leak, cheat, repeat: Data contamination and evaluation malpractices in closed-source LLMs](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*.
- William Belzak. 2020. [Testing differential item functioning in small samples](#). *Multivariate Behavioral Research*.
- Ryan Burnell, Wout Schellaert, John Burden, Tomer D. Ullman, Fernando Martinez-Plumed, Joshua B. Tenenbaum, Danaja Rutar, Lucy G. Cheke, Jascha Sohl-Dickstein, Melanie Mitchell, Douwe Kiela, Murray Shanahan, Ellen M. Voorhees, Anthony G. Cohn, Joel Z. Leibo, and Jose Hernandez-Orallo. 2023. [Rethink reporting of evaluation results in AI Science](#).
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiuyan Zhang. 2023. [Quantifying memorization across neural language models](#). In *International Conference on Learning Representations (ICLR)*.
- Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Kaijie Zhu, Hao Chen, Linyi Yang, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, Wei Ye, Yue Zhang,

801	Yi Chang, Philip S. Yu, Qiang Yang, and Xing Xie.	Zeeraq Talat, Pontus Stenetorp, and Robin Jia. 2021.	857
802	2024. A survey on evaluation of large language mod-	Dynabench: Rethinking benchmarking in NLP . In	858
803	els . <i>ACM Transactions on Intelligent Systems and</i>	<i>Proceedings of the 2021 Conference of the North</i>	859
804	<i>Technology</i> .	<i>American Chapter of the Association for Computa-</i>	860
		<i>tional Linguistics (NAACL)</i> .	861
805	Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anasta-	Seonghoon Kim and Michael Kolen. 2019. Application	862
806	sios N. Angelopoulos, Tianle Li, Dacheng Li, Hao	of IRT fixed parameter calibration to multiple-group	863
807	Zhang, Banghua Zhu, Michael I. Jordan, Joseph E.	test data . <i>Applied Measurement in Education</i> .	864
808	Gonzalez, and Ion Stoica. 2024. Chatbot arena: An		
809	open platform for evaluating LLMs by human prefer-		
810	ence . In <i>Proceedings of the International Conference</i>		
811	<i>on Machine Learning (ICML)</i> .		
812	Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot,	Alon Kipnis, Konstantinos Voudouris, Luca Maria	865
813	Ashish Sabharwal, Carissa Schoenick, and Oyvind	Schulze Buschoff, and Eric Schulz. 2024. metabench	866
814	Tafjord. 2018. Think you have solved question an-	– a sparse benchmark of reasoning and knowledge in	867
815	swering? try ARC, the AI2 reasoning challenge.	large language models . <i>Preprint</i> , arXiv:2407.12844.	868
816	<i>arXiv preprint arXiv:1803.05457</i> .		
817	Brian E. Clauser and Kathleen M. Mazor. 1998. Using	John Lalor, Hao Wu, and Hong Yu. 2019. Learning	869
818	statistical procedures to identify differentially func-	latent parameters without human response patterns:	870
819	tioning test items . <i>Educational Measurement: Issues</i>	Item response theory with artificial crowds . In <i>Pro-</i>	871
820	<i>and Practice</i> .	<i>ceedings of the 2019 Conference on Empirical Meth-</i>	872
		<i>ods in Natural Language Processing and the 9th In-</i>	873
821	Jasper Dekoninck, Mark Niklas Mueller, and Martin	<i>ternational Joint Conference on Natural Language</i>	874
822	Vechev. 2024. ConStat: Performance-based contam-	<i>Processing (EMNLP-IJCNLP)</i> .	875
823	ination detection in large language models . In <i>Ad-</i>		
824	<i>vances in Neural Information Processing Systems</i>	Peiyang Li, Xiaobo Tang, Sai Chen, Ying Cheng,	876
825	<i>(NeurIPS)</i> .	Ronald Metoyer, Tianyi Hua, and Nitesh V. Chawla.	877
		2025. ATLAS: Adaptive testing for LLM evalua-	878
826	John R. Donoghue and Nancy L. Allen. 1993. An em-	tion: A psychometric alternative to static benchmarks .	879
827	pirical examination of the Mantel-Haenszel statistic	<i>Preprint</i> , arXiv:2511.04689.	880
828	for the study of differential item functioning. <i>Applied</i>		
829	<i>Measurement in Education</i> , 6(2):163–179.	Yucheng Li, Frank Guerin, and Chenghua Lin. 2024a.	881
830	Shahriar Golchin and Mihai Surdeanu. 2024. Time	LatestEval: Addressing data contamination in lan-	882
831	travel in LLMs: Tracing data contamination in large	guage model evaluation through dynamic and time-	883
832	language models . In <i>Proceedings of the International</i>	sensitive test construction . In <i>Proceedings of the</i>	884
833	<i>Conference on Learning Representations (ICLR)</i> .	<i>AAAI Conference on Artificial Intelligence (AAAI)</i> .	885
834	Eitam Habba, Ido Itzhak, Asaf Yehudai, Yotam Per-	Yucheng Li, Yunhao Guo, Frank Guerin, and Chenghua	886
835	litiz, Elron Bandel, Michal Shmueli-Scheuer, Leshem	Lin. 2024b. An open-source data contamination re-	887
836	Choshen, and Gabriel Stanovsky. 2026. Grow-	port for large language models . In <i>Findings of the</i>	888
837	ing pains: Extensible and efficient LLM bench-	<i>Association for Computational Linguistics: EMNLP</i>	889
838	marking via fixed parameter calibration . <i>Preprint</i> ,	2024.	890
839	arXiv:2604.12843.	Rongzhen Liang. 2026. DVD: Detection via variance	891
		of generation distribution for variant contamination .	892
840	Matthew Hardy, Juan Gilbert, and Benjamin Domingue.	<i>Preprint</i> , arXiv:2601.04895.	893
841	2026. Efficient detection of bad benchmark		
842	items with novel scalability coefficients . <i>Preprint</i> ,	Stephanie Lin, Jacob Hilton, and Owain Evans. 2022.	894
843	arXiv:2603.24999.	TruthfulQA: Measuring how models mimic human	895
844	Dan Hendrycks, Collin Burns, Steven Basart, Andy	falsehoods . In <i>Proceedings of the 60th Annual Meet-</i>	896
845	Zou, Mantas Mazeika, Dawn Song, and Jacob Stein-	<i>ing of the Association for Computational Linguistics</i>	897
846	hardt. 2021. Measuring massive multitask language	<i>(ACL)</i> , pages 3214–3252.	898
847	understanding . In <i>Proceedings of the International</i>	Frederic M. Lord. 1980. <i>Applications of Item Response</i>	899
848	<i>Conference on Learning Representations (ICLR)</i> .	<i>Theory to Practical Testing Problems</i> . Lawrence	900
849	Paul W. Holland and Dorothy T. Thayer. 1988. Dif-	Erlbaum Associates, Hillsdale, NJ.	901
850	ferential item functioning and the Mantel-Haenszel		
851	procedure . In Howard Wainer and Henry I. Braun,	Zhimeng Luo, Lixin Wu, Adam Frisch, and Daqing	902
852	editors, <i>Test Validity</i> . Lawrence Erlbaum Associates.	He. 2025. Measuring competency, not performance:	903
		Item-aware evaluation across medical benchmarks .	904
853	Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh	<i>Preprint</i> , arXiv:2509.24186.	905
854	Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vid-		
855	gen, Grusha Prasad, Amanpreet Singh, Pratik Ring-	Hotaka Maeda. 2025. Finding words associated with	906
856	shia, Zhiyi Ma, Tristan Thrush, Sebastian Riedel,	DIF: Predicting DIF using LLMs and explainable AI .	907
		<i>arXiv preprint arXiv:2502.07017</i> .	908

909	Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. tinyBenchmarks: evaluating LLMs with fewer examples . In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> .	964
910		965
911		966
912		967
913		968
914	William Meredith. 1993. Measurement invariance, factor analysis and factorial invariance . <i>Psychometrika</i> .	969
915		970
916	Bengt Muthén. 1985. A method for studying the homogeneity of test items with respect to other relevant variables . <i>Journal of Educational Statistics</i> .	971
917		972
918		973
919	Ratna Nandakumar. 1993. Simultaneous DIF amplification and cancellation: Shealy-Stout’s test for DIF. <i>Journal of Educational Measurement</i> , 30(4):293–311.	974
920		975
921		976
922		977
923	Qi Qian and Chengsong Huang. 2026. Benchmark²: Systematic evaluation of LLM benchmarks . <i>Preprint</i> , arXiv:2601.03986.	978
924		979
925		
926	Nambury S. Raju. 1988. The area between two item characteristic curves . <i>Psychometrika</i> .	980
927		981
928	Pedro Rodriguez, Joe Barrow, Alexander Miserlis Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. 2021. Evaluation examples are not equally informative: How should that change NLP leaderboards? In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)</i> , pages 4486–4503.	982
929		983
930		984
931		
932		
933		
934		
935	Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2020. WinoGrande: An adversarial Winograd schema challenge at scale. In <i>Proceedings of the AAAI Conference on Artificial Intelligence</i> , volume 34, pages 8732–8740.	985
936		986
937		987
938		988
939		
940	Robin Shealy and William Stout. 1993. A model-based standardization approach that separates true bias/DIF from group ability differences and detects test bias/DTF as well as item bias/DIF . <i>Psychometrika</i> .	989
941		990
942		991
943		
944		
945	Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. 2024. Detecting pretraining data from large language models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	992
946		993
947		994
948		995
949		996
950		
951	Hariharan Swaminathan and H. Jane Rogers. 1990. Detecting differential item functioning using logistic regression procedures . <i>Journal of Educational Measurement</i> .	997
952		998
953		999
954		1000
955	Yongding Tao, Tian Wang, Yihong Dong, Huanyu Liu, Kechi Zhang, Xiaolong Hu, and Ge Li. 2026. Detecting data contamination from reinforcement learning post-training for large language models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	1001
956		1002
957		1003
958		1004
959		1005
960		1006
961	Jeanne A. Teresi and John A. Fleishman. 2007. Differential item functioning and health assessment . <i>Quality of Life Research</i> .	1007
962		1008
963		
	Kushal Tirumala, Aram Markosyan, Luke Zettlemoyer, and Armen Aghajanyan. 2022. Memorization without overfitting: Analyzing the training dynamics of large language models . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	1009
		1010
		1011
		1012
	Robert J. Vandenberg and Charles E. Lance. 2000. A review and synthesis of the measurement invariance literature: Suggestions, practices, and recommendations for organizational research . <i>Organizational Research Methods</i> .	1013
		1014
		1015
		1016
		1017
		1018
	Clara Vania, Phu Mon Htut, William Huang, Dhara Mungra, Richard Yuanzhe Pang, Jason Phang, Haokun Liu, Kyunghyun Cho, and Samuel R. Bowman. 2021. Comparing test sets with item response theory . In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	
	Haonan Wang, Haoyang Li, Boeun Ko, and Hongyang Zhang. 2026. On the fragility of benchmark contamination detection in reasoning models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	
	Colin White. 2025. LiveBench: A challenging, contamination-free LLM benchmark . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	
	Carol M. Woods. 2009. Empirical selection of anchors for tests of differential item functioning . <i>Applied Psychological Measurement</i> .	
	Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farzaneh, and Yejin Choi. 2019. HellaSwag: Can a machine really finish your sentence? In <i>Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)</i> , pages 4791–4800.	
	Jingyang Zhang, Jingwei Sun, Eric Yeats, Yang Ouyang, Martin Kuo, Jianyi Zhang, Hao Yang, and Hai Li. 2024. Min-K%++: Improved baseline for detecting pre-training data from large language models . <i>Preprint</i> , arXiv:2404.02936.	
	Hao Zhou, Heng Huang, Zhenheng Zhao, Lin Han, Haoyu Wang, Ke Chen, Min Yang, Weidong Bao, Ji Dong, Bo Xu, Changshui Zhu, Huijun Cao, and Tiejun Zhao. 2026. Lost in benchmarks? rethinking large language model benchmarking with item response theory . In <i>Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)</i> .	
	Longyuan Zhu, Hairan Hua, Linlin Miao, and Bing Zhao. 2026. Benchmark health index: A systematic framework for benchmarking the benchmarks of LLMs. <i>arXiv preprint arXiv:2602.11674</i> .	
	Bruno D. Zumbo. 1999. A Handbook on the Theory and Methods of Differential Item Functioning (DIF): Logistic Regression Modeling as a Unitary Framework for Binary and Likert-Type (Ordinal) Item Scores . Directorate of Human Resources Research and Evaluation, Department of National Defense, Ottawa.	

A Method Details

A.1 Diagnostics

Three diagnostic procedures complement the primary DIF analysis.

Local independence (Yen’s Q3). The MH procedure assumes that item responses are conditionally independent given θ . Yen’s Q3 statistic measures residual correlations between item pairs after conditioning on the latent trait. We compute Q3 for all within-subject and cross-subject item pairs, with $|Q3| > 0.20$ as the threshold for local dependence. Elevated within-subject dependence would indicate that items within the same subject share variance beyond what θ captures, which could inflate within-subject DIF rates.

Domain-specificity. We test whether temporal DIF concentrates in domains where models improved most. At the ecological level, we compute Spearman ρ between subject-level C% and subject-level accuracy improvement from 2023 to 2024. At the item level, logistic regression predicts item-level DIF C classification from domain-level accuracy change, with and without item property controls (difficulty, discrimination). This decomposition quantifies how much of the observed instability is attributable to aggregate domain-level gains versus item-specific behavioral change.

Instability decomposition. When domain accuracy improvement is included alongside item-level controls, we compute the counterfactual C% that would remain after removing domain-level effects. The gap between observed and counterfactual C% measures the contribution of domain-level improvement to the total instability.

A.2 Statistical Considerations

Multiple testing. The ETS C classification imposes a dual threshold (effect size $|\Delta_{MH}| \geq 1.5$ and significance $p < 0.05$), which provides implicit control over false positives. As a robustness check, we apply Benjamini-Hochberg FDR correction at $q = 0.05$. The random-split sanity check provides an empirical estimate of the false positive rate under the null hypothesis.

Threshold sensitivity. We report results across a range of DIF thresholds ($|\Delta_{MH}| \in \{1.0, 1.5, 2.0, 2.5\}$) to verify that qualitative conclusions do not depend on the specific ETS C cutoff. The standard threshold of

$|\Delta_{MH}| \geq 1.5$ corresponds to the ETS C definition used throughout the paper.

Architecture heterogeneity. The METABENCH MMLU population consists of 99.98% decoder-only models (5,226 of 5,227). This near-complete architectural homogeneity eliminates architecture as a confound for DIF analysis. We verify this by excluding the single encoder-decoder model and confirming identical results.

B Structural Analysis Details

Three progressively refined approaches fail to reduce the temporal false positive rate below 31%, revealing a structural property of the data rather than a methodological limitation. Matched-ability quintile analysis restricts comparisons to models of similar total score and yields enrichment ratios of 0.94–1.04, with temporal FPR unchanged at approximately 0.31. IRT θ -conditioning with 20 ability strata achieves $r(\theta, \text{total score}) = 0.9986$ and $\text{FPR} = 0.312$, confirming that the external total score already captures nearly all ability variance relevant to DIF computation. Within-subject DIF analysis on 5 representative subjects, intended to isolate subject-specific effects, finds mean C% = 38.1%, exceeding the 25% threshold that would indicate viable bifactor structure. Table 3 summarizes these results.

The temporal FPR reflects genuine item-level instability that persists regardless of how ability is measured or conditioned. Yen’s Q3 diagnostic reveals mild local dependence within subjects (11.3% of within-subject item pairs exceed $|Q3| > 0.20$, vs. 2.2% cross-subject; $5\times$ ratio, $p = 2.3 \times 10^{-9}$), which contributes to elevated within-subject DIF rates but is not the primary driver: the cross-subject matching protocol addresses this bias, and temporal C% persists under external matching (full Q3 distributions in Appendix Figure A2).

Table 3: Structural approaches to reducing temporal DIF. None reduces baseline FPR below 31%, indicating that the observed instability is not an artifact of ability matching. Subscript $\pm$ values denote standard deviations across 5 injection seeds (baseline row only).

Approach	Mechanism	Temp. FPR	TPR ($p=0.5$)	Outcome
Baseline (ext.)	Cross-bench, total score	0.314 $\pm$.009	0.834 $\pm$.027	Reference
Matched-ability	Within-quintile temp. DIF	~ 0.31	–	Enrich. 0.94–1.04
θ -cond.	θ replaces total score	0.312	0.876	$r(\theta, \text{total})=0.9986$
Within-subject	Per-subject DIF	–	–	C% = 38.1% (>25%)

C Additional Figures




Figure 3: Cross-subject vs. within-subject matching for semi-synthetic DIF detection. Cross-subject matching eliminates FPR inflation ($\Delta\text{FPR} = 0.000$) while preserving detection power ($\Delta\text{TPR} > 0.3$ for 4/5 subjects at $p = 0.5$). Within-subject matching inflates FPR by $+0.186$ on average.




Figure 4: Downstream impact of DIF-based item filtering on model rankings. Despite $\rho \approx 0.997$ overall, individual model families show systematic rank shifts when evaluated on stable items only.

D Cross-Benchmark Comparison




Figure 5: GSM1K contamination gap vs. DIF ranking advantage ($N = 33$). The negative correlation ($r = -0.478$, $p = 0.005$) indicates that models with larger contamination gaps benefit less from DIF-flagged items.

Table 4: Cross-benchmark comparison of DIF analysis on MMLU and GSM8K. Both benchmarks exhibit temporal measurement drift with validated semi-synthetic detection power. Standard and external matching use different exploitation rates (p) because each benchmark’s optimal operating point differs; external matching results are compared at $p = 0.5$ for both. Subscript $\pm$ values denote standard deviations across 5 random injection seeds.

Metric	MMLU	GSM8K
Items	12,508	1,319
Models	5,227	6,702
Temporal C%	31.3%	18.9%
Sanity C%	0.18%	0.3%
ΔTPR ($p=0.7$, standard)	—	$+0.818 \pm .006$
ΔTPR ($p=0.5$, external)	$+0.508 \pm .027$	$+0.514 \pm .028$